

Author: Onyedikachukwu Okonkwo Project: Lloyd Bank EDA Project

```
# install libraries
!python -m pip install --upgrade pip -q
!pip install pandas numpy matplotlib seaborn scikit-learn openpyxl -q
!pip install notebook -q
!pip install nbconvert -q
!pip install tabulate -q
```

Lloyd Bank EDA

Bottom Line Up Front

- Discovered through clustering that our **Most** valuable customers in Group 6 and 8 churn with a Churn Rate at rates higher than average.
- Group 6 spend 0.44 above average but churn at 0.21
- Group 8 spend 0.36 above average but churn at 0.27
- They are our most valuable customers and we cannot figure out why they are leaving, because the features we rely on have such poor correlation a coin flip is better at predicting churn.
- The group with the lowest churn is our **Least** Valuable customers. They spend the least.
- The data is all over the place, and no single feature seems to correlate to the why they churn at such a high clip. Further analysis is required.

Project Overview

Objective: First Critical Steps to build a predictive model for customer churn. Gather all relevant data, and then conduct an EDA and prepare the data set for model development. Key Results:

Business Problem: We need to predict which of our customers are likely to churn.

Goal: Conduct exploratory data analysis and translate findings into a predictive model that can predict which customers are going to churn. Methodology and Justification.

Tools Used: Python for data cleaning, manipulation, analysis, and visualization Reason: Python has libraries that are suited for the assignment. Pandas is used to manipulate the data, numpy for calculation, matplotlib and seaborn for data visualization, sklearn for machine learning model, and datetime to calculate date ranges in days.

```
#import libraries

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn
from datetime import datetime
import warnings
warnings.filterwarnings('ignore')
```

Identify and Gathering Steps:

- Downloaded the data from the excel files, and uploaded a copy using pd.read_excel as it is an excel file.
- Kept the names the same as the sheets for consistency across the board.
- Printed the head of each Sheet and compared to original dataset for accuracy and to ensure consistency and nothing is lost

```
# Set Display Options
pd.set_option('display.max_columns', None) # Set option to display all columns
pd.set_option('display.float_format', '{:.2f}'.format) # Set float format to 2 decimal places

Churn_Status = pd.read_excel('Customer_Churn_Data_Large.xlsx',
sheet_name='Churn_Status')
Customer_Demographics =
pd.read_excel('Customer_Churn_Data_Large.xlsx',
sheet_name='Customer_Demographics')
Transaction_History = pd.read_excel('Customer_Churn_Data_Large.xlsx',
sheet_name='Transaction_History')
Customer_Service = pd.read_excel('Customer_Churn_Data_Large.xlsx',
sheet_name='Customer_Service')
Online_Activity = pd.read_excel('Customer_Churn_Data_Large.xlsx',
sheet_name='Online_Activity')

# Display first few rows of each dataset
print("Churn Status:")
print(Churn_Status.head())
print("\nCustomer Demographics:")
print(Customer_Demographics.head())
print("\nTransaction History:")
print(Transaction_History.head())
print("\nCustomer Service:")
print(Customer_Service.head())
print("\nOnline Activity:")
print(Online_Activity.head())
```

```
# Display info of each dataset to understand structure and data types
print(Customer_Demographics.info())
print(Transaction_History.info())
print(Churn_Status.info())
print(Customer_Service.info())
print(Online_Activity.info())
```

Churn Status:

	CustomerID	ChurnStatus
0	1	0
1	2	1
2	3	0
3	4	0
4	5	0

Customer Demographics:

	CustomerID	Age	Gender	MaritalStatus	IncomeLevel
0	1	62	M	Single	Low
1	2	65	M	Married	Low
2	3	18	M	Single	Low
3	4	21	M	Widowed	Low
4	5	21	M	Divorced	Medium

Transaction History:

	CustomerID	TransactionID	TransactionDate	AmountSpent
0	1	7194	2022-03-27	416.50
Electronics				
1	2	7250	2022-08-08	54.96
Clothing				
2	2	9660	2022-07-25	197.50
Electronics				
3	2	2998	2022-01-25	101.31
Furniture				
4	2	1228	2022-07-24	397.37
Clothing				

Customer Service:

	CustomerID	InteractionID	InteractionDate	InteractionType
0	1	6363	2022-03-31	Inquiry
Resolved				
1	2	3329	2022-03-17	Inquiry
Resolved				
2	3	9976	2022-08-24	Inquiry
Resolved				
3	4	7354	2022-11-18	Inquiry
Resolved				
4	4	5393	2022-07-03	Inquiry

Unresolved

Online Activity:

```
CustomerID  LastLoginDate  LoginFrequency  ServiceUsage
0           1            2023-10-21          34   Mobile App
1           2            2023-12-05          5    Website
2           3            2023-11-15          3    Website
3           4            2023-08-25          2    Website
4           5            2023-10-27          41   Website
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1000 entries, 0 to 999

Data columns (total 5 columns):

#	Column	Non-Null Count	Dtype
0	CustomerID	1000 non-null	int64
1	Age	1000 non-null	int64
2	Gender	1000 non-null	object
3	MaritalStatus	1000 non-null	object
4	IncomeLevel	1000 non-null	object

dtypes: int64(2), object(3)

memory usage: 39.2+ KB

None

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 5054 entries, 0 to 5053

Data columns (total 5 columns):

#	Column	Non-Null Count	Dtype
0	CustomerID	5054 non-null	int64
1	TransactionID	5054 non-null	int64
2	TransactionDate	5054 non-null	datetime64[ns]
3	AmountSpent	5054 non-null	float64
4	ProductCategory	5054 non-null	object

dtypes: datetime64[ns](1), float64(1), int64(2), object(1)

memory usage: 197.6+ KB

None

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1000 entries, 0 to 999

Data columns (total 2 columns):

#	Column	Non-Null Count	Dtype
0	CustomerID	1000 non-null	int64
1	ChurnStatus	1000 non-null	int64

dtypes: int64(2)

memory usage: 15.8 KB

None

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1002 entries, 0 to 1001

Data columns (total 5 columns):

#	Column	Non-Null Count	Dtype
---	--------	----------------	-------

```

-----|-----|-----|-----|
0   CustomerID      1002 non-null  int64
1   InteractionID  1002 non-null  int64
2   InteractionDate 1002 non-null  datetime64[ns]
3   InteractionType 1002 non-null  object
4   ResolutionStatus 1002 non-null  object
dtypes: datetime64[ns](1), int64(2), object(2)
memory usage: 39.3+ KB
None
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 4 columns):
 #   Column           Non-Null Count  Dtype  
-----|-----|-----|-----|
 0   CustomerID      1000 non-null   int64  
 1   LastLoginDate   1000 non-null   datetime64[ns]
 2   LoginFrequency  1000 non-null   int64  
 3   ServiceUsage    1000 non-null   object 
dtypes: datetime64[ns](1), int64(2), object(1)
memory usage: 31.4+ KB
None

```

Data Cleaning Process

- Created a backup of the original raw data
- Created a Transaction_Summary based on Transaction_History grouped by CustomerID to create a single column per customer rather than multiple columns.
- Grouped by minimum and Maximum transaction, the amount of transactions, and the total spent per customers, also calculated the LoyaltyLength by subtracting the most recent transaction date by the first transaction date.

All these are new columns that are added to a table and dropped all null values from Transaction_Summary I don't need to analyze customers who haven't spend any money.

Documentation: All cleaning steps are documented with inline comments.

- For the Customer_Service, and InteractionType, I created unique columns again to make sure there is one column per CustomerID as it has a one to many relationship with InteractionType. 1 customer can put in multiple Interaction or have none, The summary dataframe groups it by feedback, inquiry or complaint and if it was resolved or not each has its own columns assigned to that customer.
- For customers with no values the Nan are filled with 0 rather than dropping the column
- For customer with no demographics or churn status or online activity, they Nan Values were dropped.
- The result was we started with 1000 rows and ended up with 1000 rows.
- Customer_Data is a new table that groups all these tables and columns together using a merge and inner join for the Churn_Status and the Demographics on the CustomerID because they are the key place holders. Then a left join with Transaction_Summary, Customer_Service_Summary, and the Online_Activity.

- Consolidated all the data into Customer_Data and encoded the non_numeric columns to prepare for learning model.

Data Quality Assessment (ROCCC):

- Reliable: Direct from mock company data
- Original: First-party company data
- Comprehensive: Contains all necessary variables for analysis
- Current: Static dataset (not real-time)
- Cited: Well-documented source

Limitations:

- Data does not indicate if feedback is positive or negative which is a valid metric in measuring which customers churn.
- Data is not dynamic/real-time
- Limited to a point-in-time snapshot

```
# Create New Columns based on Transaction History of each "CustomerID"
Transaction_Summary = Transaction_History.groupby('CustomerID').agg(
    TotalSpent = ("AmountSpent", "sum"), # Create column for total
    transaction amount
    MinTransaction = ("AmountSpent", "min"), # Create column for
    minimum transaction amount
    MaxTransaction = ("AmountSpent", "max"), # Create column for
    maximum transaction amount
    TransactionFrequency = ("AmountSpent", "count"), # Create column
    for frequency of transactions
    MostRecentTransaction = ("TransactionDate", "max"), # Create
    column for most recent transaction date
    FirstTransaction = ("TransactionDate", "min") # Create column for
    first transaction date
).reset_index()

# Create column for length of relationship based on first transaction
and most recent transaction
Transaction_Summary['LoyaltyLength'] =
(Transaction_Summary['MostRecentTransaction'] -
Transaction_Summary['FirstTransaction']).dt.days
Transaction_Summary = Transaction_Summary.fillna(0) # Fill any NaN
values with 0

# Check to make sure all is right with the cosmos thus far
display(Transaction_Summary.head())
```

	CustomerID	TotalSpent	MinTransaction	MaxTransaction	\
0	1	416.50	416.50	416.50	
1	2	1547.42	54.96	397.37	
2	3	1702.98	51.07	419.95	
3	4	917.29	44.22	382.39	
4	5	2001.49	69.86	475.69	

LoyaltyLength	TransactionFrequency	MostRecentTransaction	FirstTransaction
0	1	2022-03-27	2022-03-27
1	7	2022-11-19	2022-01-09
2	6	2022-10-08	2022-02-11
3	5	2022-12-27	2022-05-22
4	8	2022-12-21	2022-02-21
303			

```

Interaction_Types = Customer_Service['InteractionType'].unique() # 
Make a List of the different interaction types

summaries = [] # Create an empty list to hold the summaries of each
interaction type

for interaction_type in Interaction_Types: # for loop that loops
through each interaction type and creates a summary dataframe
    type_summary =
Customer_Service[Customer_Service['InteractionType'] ==
interaction_type].groupby('CustomerID').agg(**{ # group by CustomerID
and aggregate the data
        f'{interaction_type}Count': ('InteractionType', 'count'), # f
string to create dynamic column names
        f'{interaction_type}Resolved': ("ResolutionStatus", lambda x:
(x == 'Resolved').sum())
    }).reset_index()

    type_summary[f'{interaction_type}Unresolved'] =
type_summary[f'{interaction_type}Count'] -
type_summary[f'{interaction_type}Resolved'] # since we cant add or
subtract in the agg function, we create a new column for unresolved
interactions by subtracting resolved from total count

    summaries.append(type_summary) # append the summary dataframe to
the list

# Merge all interaction summaries into a single dataframe
Customer_Service_Summary = summaries[0] # start with the first summary
for i in range(1, len(summaries)): # loop through the rest of the
summaries
    Customer_Service_Summary = pd.merge(Customer_Service_Summary,
summaries[i], on='CustomerID', how='outer') # merge on CustomerID with
an outer join to keep all customers

#fill NaN values with 0s for customers who did not have certain

```

```

interaction types
Customer_Service_Summary = Customer_Service_Summary.fillna(0)
# Check to make sure all is right with the cosmos thus far
display(Customer_Service_Summary.head())

   CustomerID  InquiryCount  InquiryResolved  InquiryUnresolved \
0             1          1.00            1.00            0.00
1             2          1.00            1.00            0.00
2             3          1.00            1.00            0.00
3             4          2.00            1.00            1.00
4             6          0.00            0.00            0.00

   FeedbackCount  FeedbackResolved  FeedbackUnresolved  ComplaintCount \
\
0           0.00            0.00            0.00            0.00
1           0.00            0.00            0.00            0.00
2           0.00            0.00            0.00            0.00
3           0.00            0.00            0.00            0.00
4           1.00            1.00            0.00            0.00

   ComplaintResolved  ComplaintUnresolved
0           0.00            0.00
1           0.00            0.00
2           0.00            0.00
3           0.00            0.00
4           0.00            0.00

# Address Nan values in individual datasets before merging
Customer_Demographics = Customer_Demographics.dropna()
Churn_Status = Churn_Status.dropna()
Online_Activity = Online_Activity.dropna()

# Merge Customer Demographics, Transaction Summary, Customer Service
# Summary, Online Activity, and Churn Status into a single dataframe
Customer_Data = pd.merge(Customer_Demographics, Churn_Status,
on='CustomerID', how = 'inner')
Customer_Data = pd.merge(Customer_Data, Transaction_Summary,
on='CustomerID', how = 'left').fillna(0) # Fill NaN values with 0 for
# customers with no transactions
Customer_Data = pd.merge(Customer_Data, Customer_Service_Summary,
on='CustomerID', how = 'left').fillna(0) # Fill NaN values with 0 for
# customers with no customer service interactions
Customer_Data = pd.merge(Customer_Data, Online_Activity,
on='CustomerID', how = 'left')

```

Analysis Steps:

1. **Descriptive Statistics**

- Used the describe().T function to calculate mean, median, standard deviation, minimum, 25% quartile, median 75% quartile, maximum, and then the skew and kurtosis as well for all numeric variables.
- Calculated skew and kurtosis to determine distribution shapes, and see where the outliers mostly are, and to determine where to use StandardScale or RobustScale when calculating z scores.
- Calculated the churn_rate to check for balance.

1. **Relationship Analysis**

- Correlation analysis between variables and churn_rate. -Visualization with bar graph and box plots to confirm correlation calculations, check for outliers and confirm statistical information.

```
#Create a table to summarize numerical features
description =
Customer_Data.select_dtypes(include='number').describe().T
description['skew'] =
Customer_Data.select_dtypes(include='number').skew()
description['kurtosis'] =
Customer_Data.select_dtypes(include='number').kurtosis()
description.round(2)

# Final check to make sure all is right with the cosmos thus far
print(Customer_Data.info())
display(Customer_Data.head())
display(description)

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 25 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   CustomerID      1000 non-null    int64  
 1   Age              1000 non-null    int64  
 2   Gender           1000 non-null    object  
 3   MaritalStatus    1000 non-null    object  
 4   IncomeLevel      1000 non-null    object  
 5   ChurnStatus      1000 non-null    int64  
 6   TotalSpent       1000 non-null    float64 
 7   MinTransaction   1000 non-null    float64 
 8   MaxTransaction   1000 non-null    float64 
 9   TransactionFrequency  1000 non-null  int64  
 10  MostRecentTransaction 1000 non-null  datetime64[ns]
 11  FirstTransaction  1000 non-null    datetime64[ns]
 12  LoyaltyLength     1000 non-null    int64  
 13  InquiryCount      1000 non-null    float64 
 14  InquiryResolved   1000 non-null    float64
```

```

15 InquiryUnresolved      1000 non-null   float64
16 FeedbackCount          1000 non-null   float64
17 FeedbackResolved        1000 non-null   float64
18 FeedbackUnresolved      1000 non-null   float64
19 ComplaintCount          1000 non-null   float64
20 ComplaintResolved       1000 non-null   float64
21 ComplaintUnresolved     1000 non-null   float64
22 LastLoginDate           1000 non-null   datetime64[ns]
23 LoginFrequency           1000 non-null   int64
24 ServiceUsage             1000 non-null   object
dtypes: datetime64[ns](3), float64(12), int64(6), object(4)
memory usage: 195.4+ KB
None

```

	CustomerID	Age	Gender	MaritalStatus	IncomeLevel	ChurnStatus	
TotalSpent \	0	1	62	M	Single	Low	0
416.50	1	2	65	M	Married	Low	1
1547.42	2	3	18	M	Single	Low	0
1702.98	3	4	21	M	Widowed	Low	0
917.29	4	5	21	M	Divorced	Medium	0
2001.49							

	MinTransaction	MaxTransaction	TransactionFrequency	
MostRecentTransaction \	0	416.50	416.50	1
2022-03-27	1	54.96	397.37	7
2022-11-19	2	51.07	419.95	6
2022-10-08	3	44.22	382.39	5
2022-12-27	4	69.86	475.69	8
2022-12-21				

	FirstTransaction	LoyaltyLength	InquiryCount	InquiryResolved \
0	2022-03-27	0	1.00	1.00
1	2022-01-09	314	1.00	1.00
2	2022-02-11	239	1.00	1.00
3	2022-05-22	219	2.00	1.00
4	2022-02-21	303	0.00	0.00

	InquiryUnresolved	FeedbackCount	FeedbackResolved
FeedbackUnresolved \	0	0.00	0.00

0.00							
1	0.00	0.00	0.00	0.00			
0.00							
2	0.00	0.00	0.00	0.00			
0.00							
3	1.00	0.00	0.00	0.00			
0.00							
4	0.00	0.00	0.00	0.00			
0.00							
ComplaintCount	ComplaintResolved	ComplaintUnresolved	LastLoginDate	\			
0	0.00	0.00	0.00	2023-10-21	2023-10-		
1	0.00	0.00	0.00	05	2023-12-		
2	0.00	0.00	0.00	15	2023-11-		
3	0.00	0.00	0.00	25	2023-08-		
4	0.00	0.00	0.00	27	2023-10-		
LoginFrequency	ServiceUsage						
0	34	Mobile App					
1	5	Website					
2	3	Website					
3	2	Website					
4	41	Website					
75% \		count	mean	std	min	25%	50%
CustomerID	750.25	1000.00	500.50	288.82	1.00	250.75	500.50
Age	56.00	1000.00	43.27	15.24	18.00	30.00	43.00
ChurnStatus	0.00	1000.00	0.20	0.40	0.00	0.00	0.00
TotalSpent	1791.90	1000.00	1267.07	738.59	9.80	626.68	1232.88
MinTransaction	146.40	1000.00	107.07	101.08	5.18	32.80	74.62
MaxTransaction	468.50	1000.00	390.18	107.38	9.80	342.86	429.93
TransactionFrequency	7.00	1000.00	5.05	2.60	1.00	3.00	5.00
LoyaltyLength	294.25	1000.00	208.11	109.96	0.00	137.75	240.00
InquiryCount	1.00	1000.00	0.31	0.52	0.00	0.00	0.00

InquiryResolved	1000.00	0.17	0.40	0.00	0.00	0.00
0.00						
InquiryUnresolved	1000.00	0.14	0.36	0.00	0.00	0.00
0.00						
FeedbackCount	1000.00	0.36	0.58	0.00	0.00	0.00
1.00						
FeedbackResolved	1000.00	0.20	0.43	0.00	0.00	0.00
0.00						
FeedbackUnresolved	1000.00	0.16	0.39	0.00	0.00	0.00
0.00						
ComplaintCount	1000.00	0.34	0.55	0.00	0.00	0.00
1.00						
ComplaintResolved	1000.00	0.16	0.38	0.00	0.00	0.00
0.00						
ComplaintUnresolved	1000.00	0.18	0.41	0.00	0.00	0.00
0.00						
LoginFrequency	1000.00	25.91	14.06	1.00	13.75	27.00
38.00						

	max	skew	kurtosis
CustomerID	1000.00	0.00	-1.20
Age	69.00	0.01	-1.21
ChurnStatus	1.00	1.47	0.17
TotalSpent	3386.04	0.27	-0.79
MinTransaction	496.99	1.60	2.39
MaxTransaction	499.86	-1.51	1.90
TransactionFrequency	9.00	-0.06	-1.24
LoyaltyLength	361.00	-0.73	-0.70
InquiryCount	2.00	1.43	1.09
InquiryResolved	2.00	2.20	4.01
InquiryUnresolved	2.00	2.35	4.41
FeedbackCount	2.00	1.38	0.88
FeedbackResolved	2.00	1.93	2.79
FeedbackUnresolved	2.00	2.33	4.71
ComplaintCount	2.00	1.38	0.95
ComplaintResolved	2.00	2.31	4.55
ComplaintUnresolved	2.00	2.14	3.79
LoginFrequency	49.00	-0.13	-1.18

```
# Count how many customers churned vs. did not churn
churn_counts = Customer_Data['ChurnStatus'].value_counts()
print(churn_counts)
# Calculate churn rate
churn_rate = (churn_counts[1] / churn_counts.sum()) * 100 # Percentage of customers who churned
print(f"Churn Rate: {churn_rate:.2f}%")
```

ChurnStatus	
0	796
1	204

```

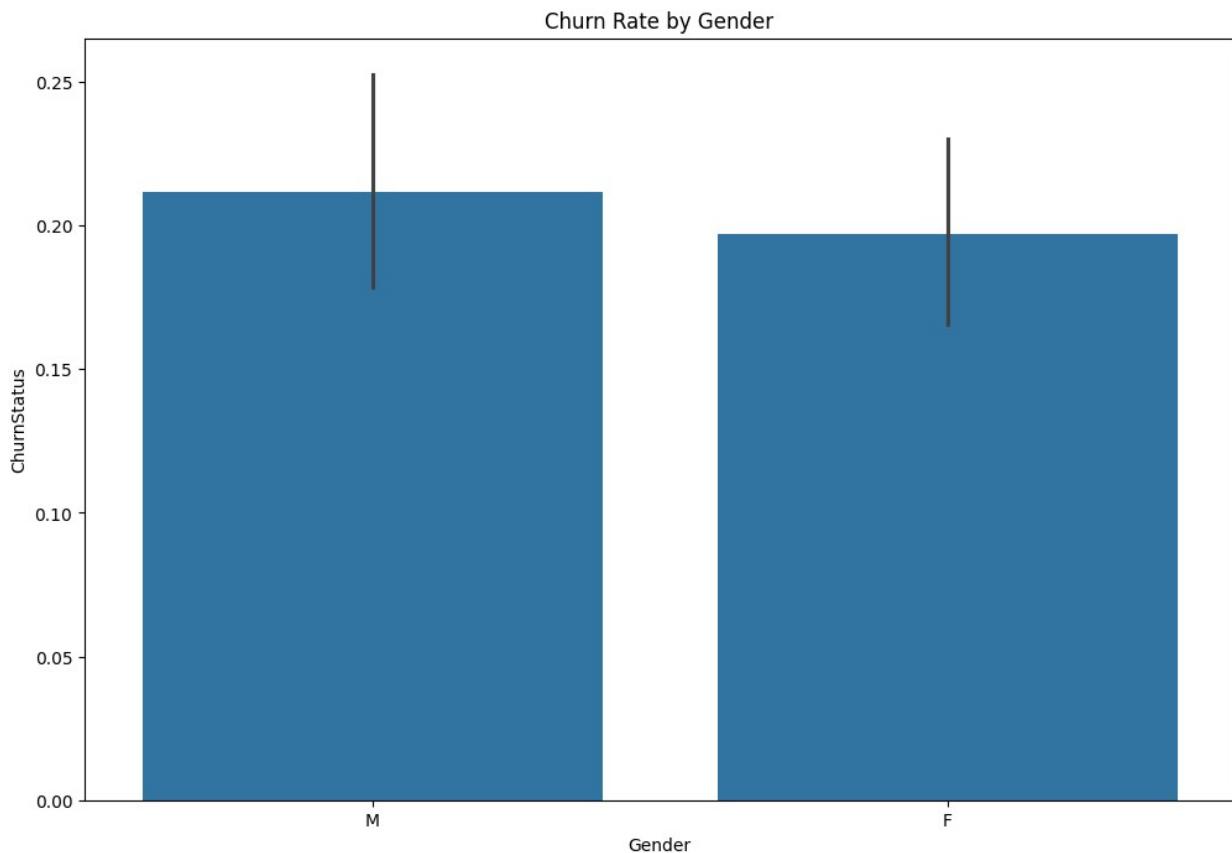
Name: count, dtype: int64
Churn Rate: 20.40%

# For Loop to calculate churn rates for categorical variables
categorical_columns =
Customer_Data.select_dtypes(include='object').columns

for col in categorical_columns:
    churn_rates = Customer_Data.groupby(col)[ 'ChurnStatus' ].mean()
    print(churn_rates)
    plt.figure(figsize=(12, 8))
    sns.barplot(x=col, y='ChurnStatus', data=Customer_Data)
    plt.title(f'Churn Rate by {col}')
    plt.show()

Gender
F   0.20
M   0.21
Name: ChurnStatus, dtype: float64

```

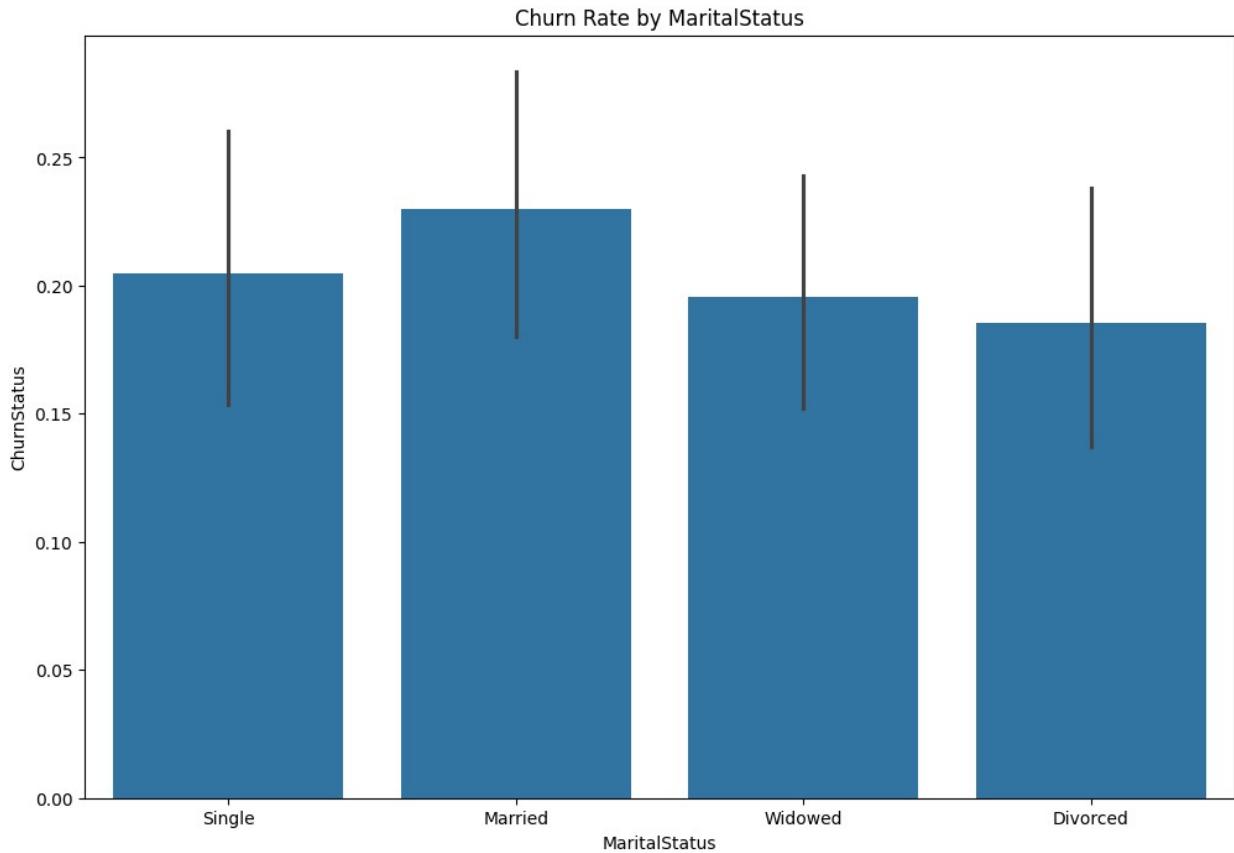


```

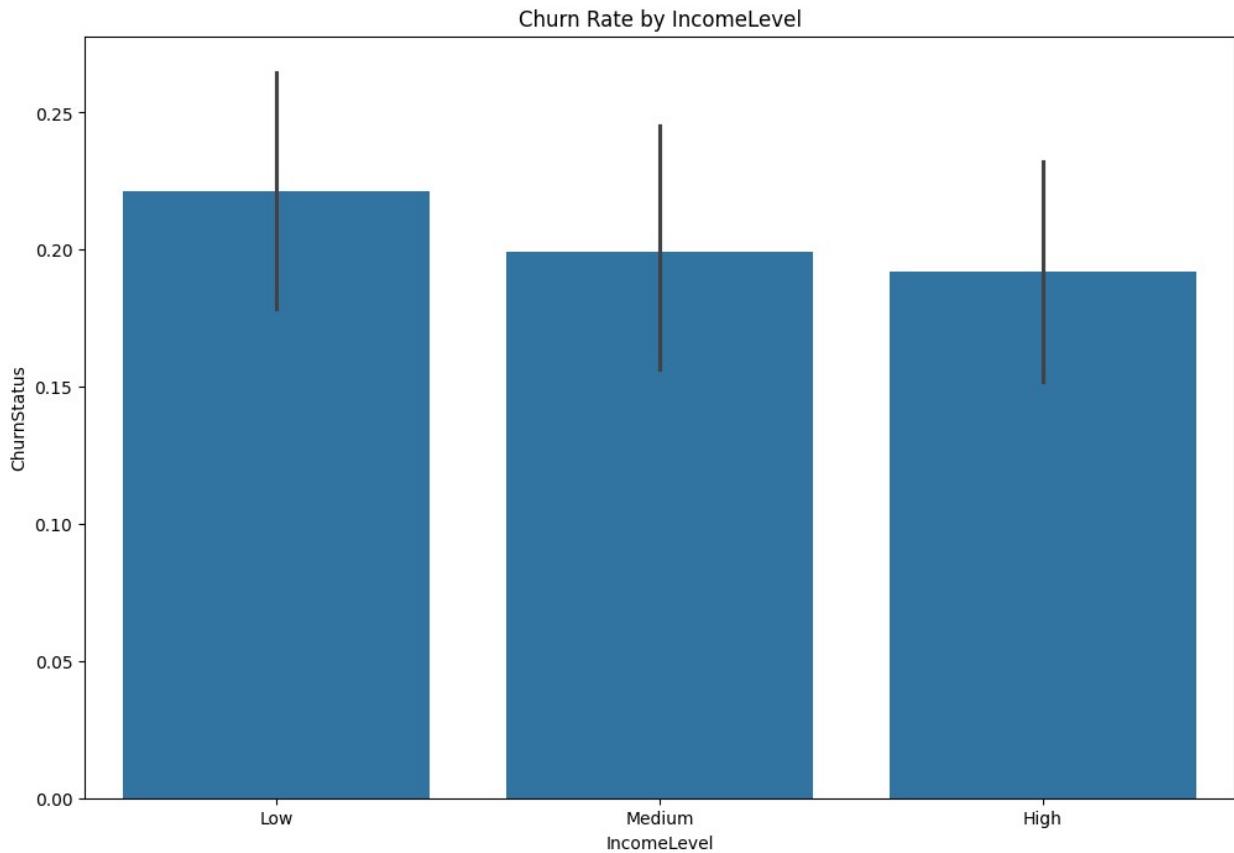
MaritalStatus
Divorced  0.19
Married   0.23

```

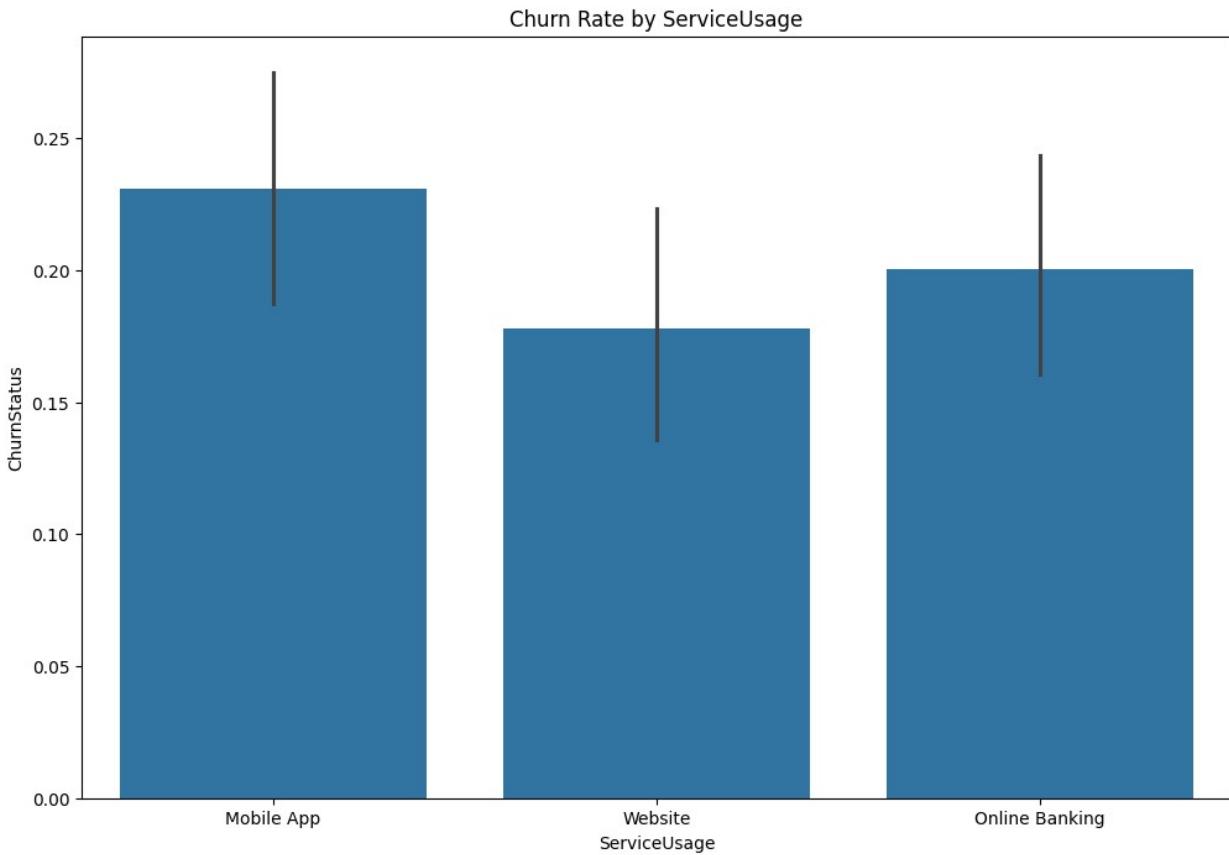
```
Single      0.20
Widowed     0.20
Name: ChurnStatus, dtype: float64
```



```
IncomeLevel
High      0.19
Low       0.22
Medium    0.20
Name: ChurnStatus, dtype: float64
```

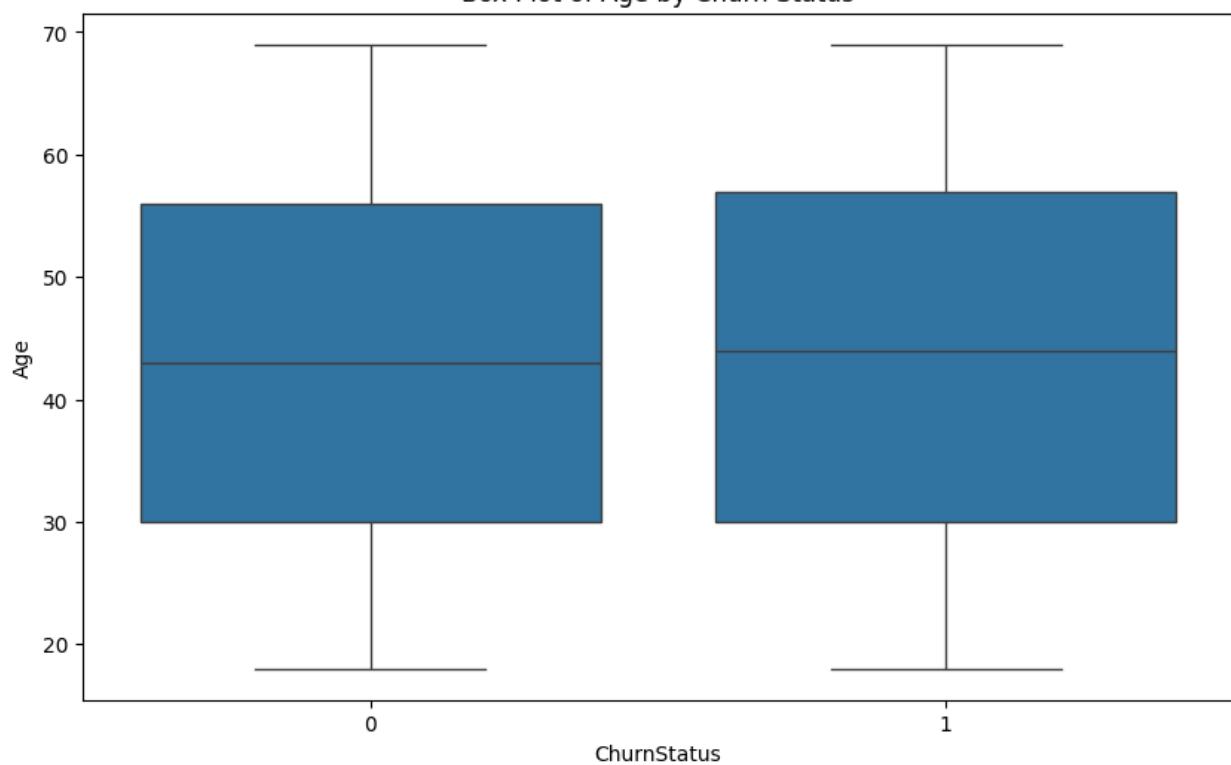


```
ServiceUsage
Mobile App      0.23
Online Banking   0.20
Website          0.18
Name: ChurnStatus, dtype: float64
```

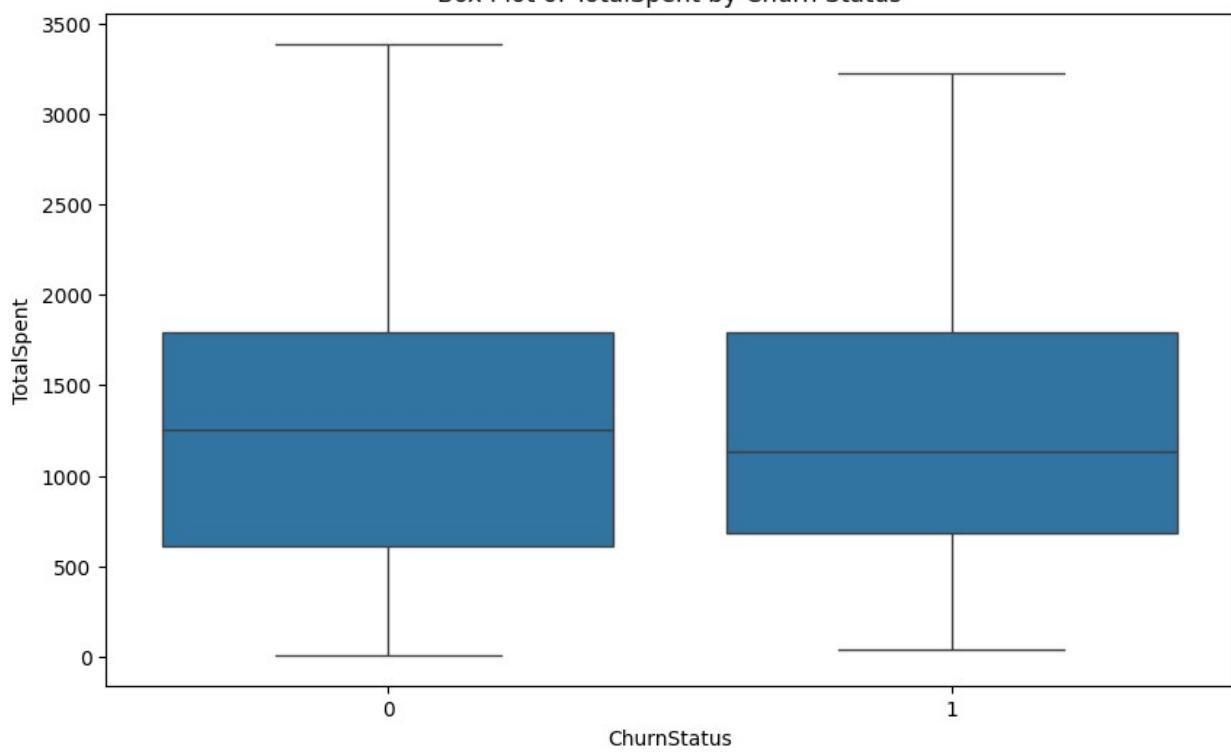


```
# For Loop to plot box plots for the numerical variables against churn status
numerical_columns =
Customer_Data.select_dtypes(include='number').columns.drop(['ChurnStatus','CustomerID'])
for col in numerical_columns:
    plt.figure(figsize=(10, 6))
    sns.boxplot(x='ChurnStatus', y=col, data=Customer_Data)
    plt.title(f'Box Plot of {col} by Churn Status')
    plt.show()
```

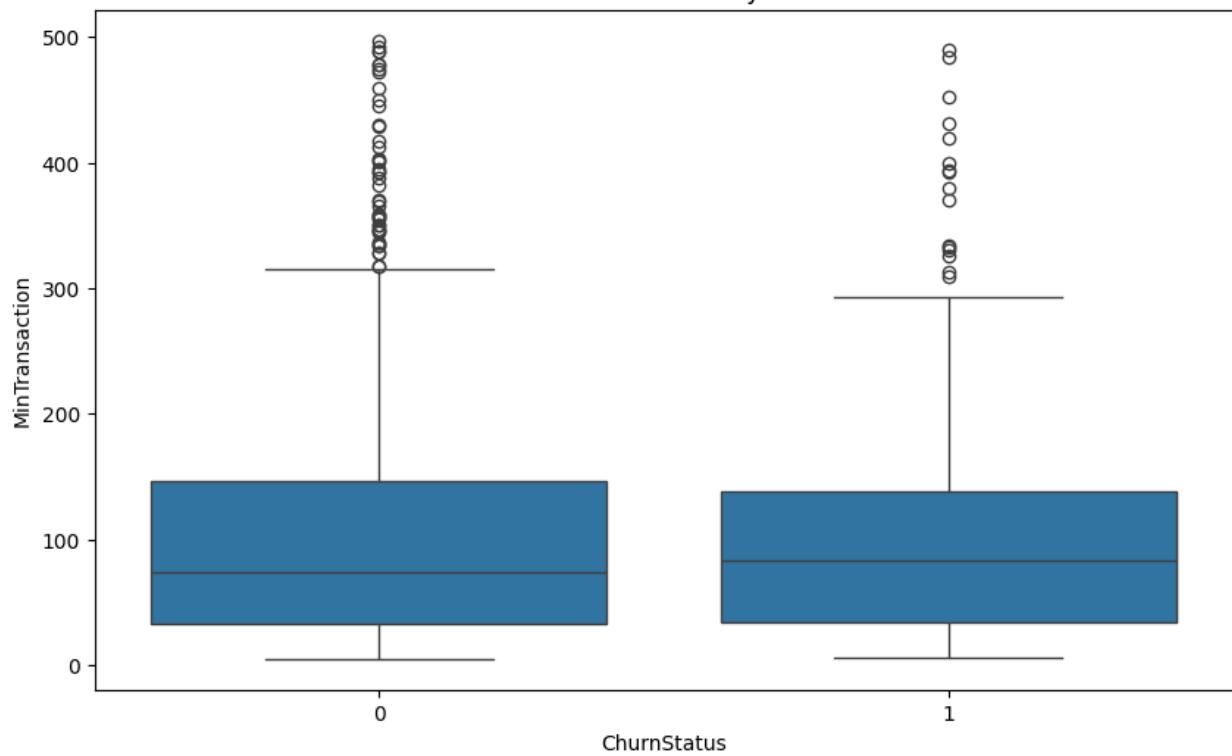
Box Plot of Age by Churn Status



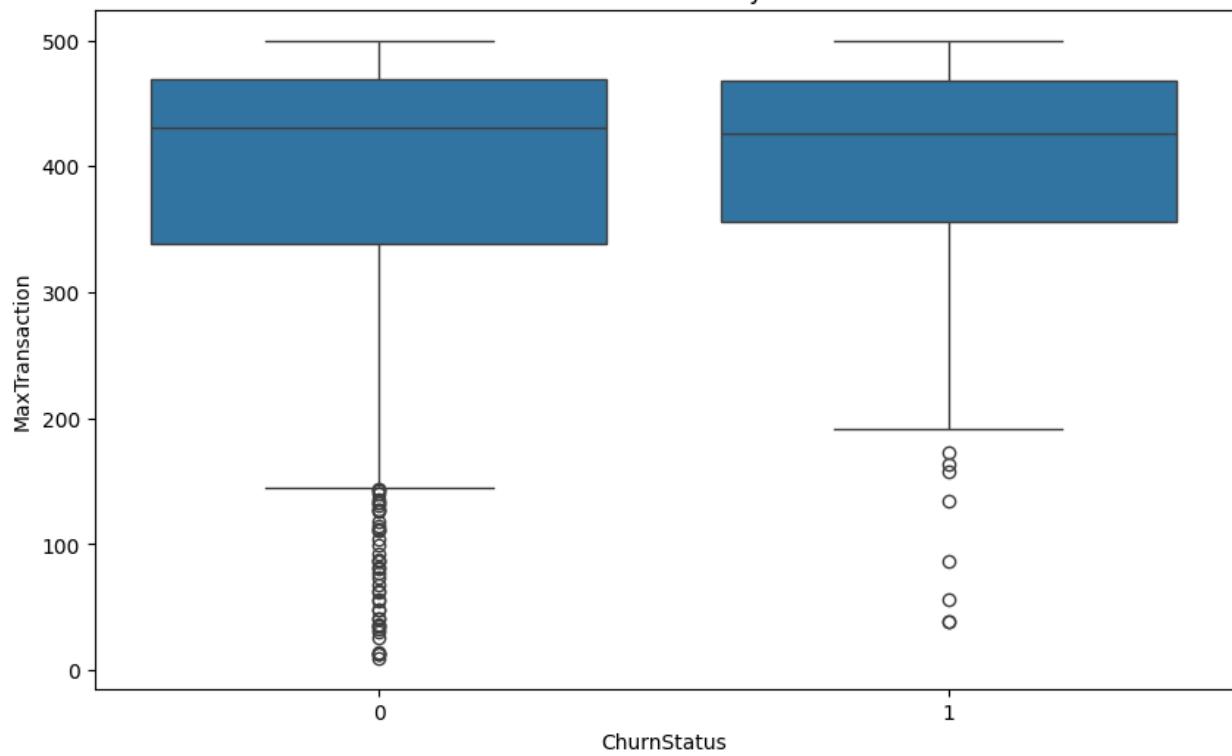
Box Plot of TotalSpent by Churn Status



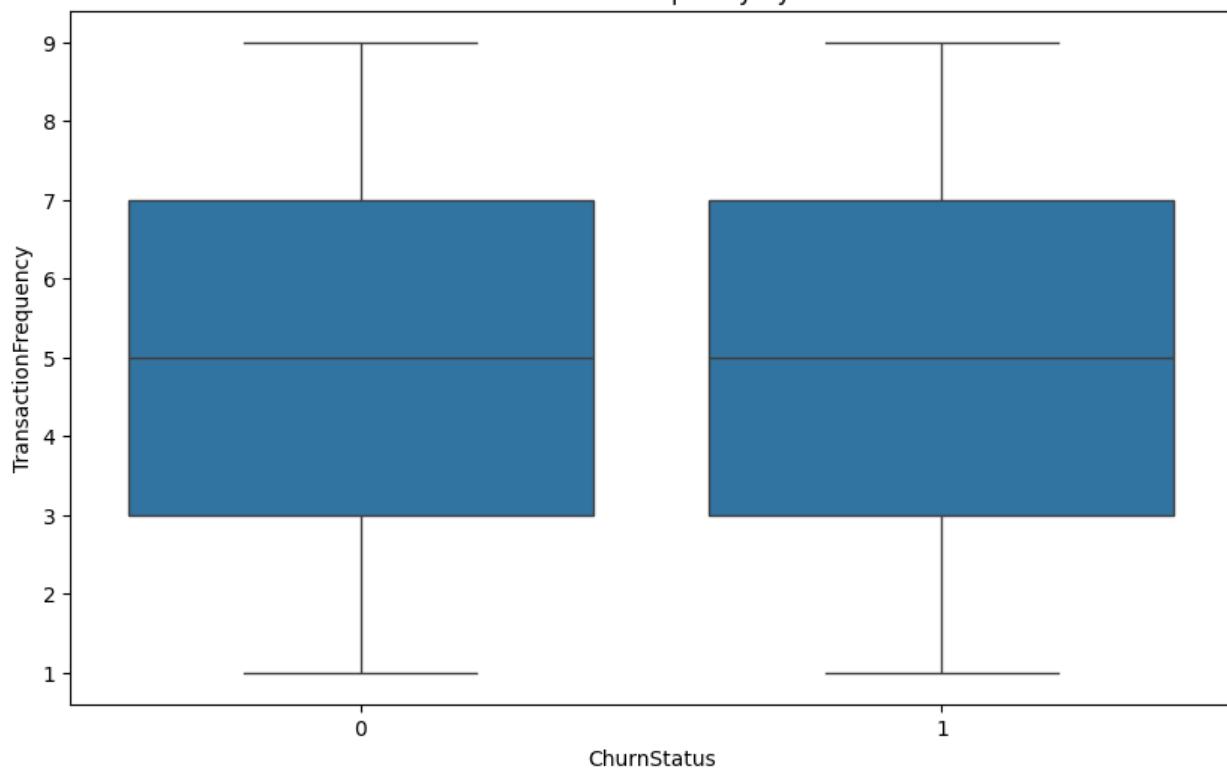
Box Plot of MinTransaction by Churn Status



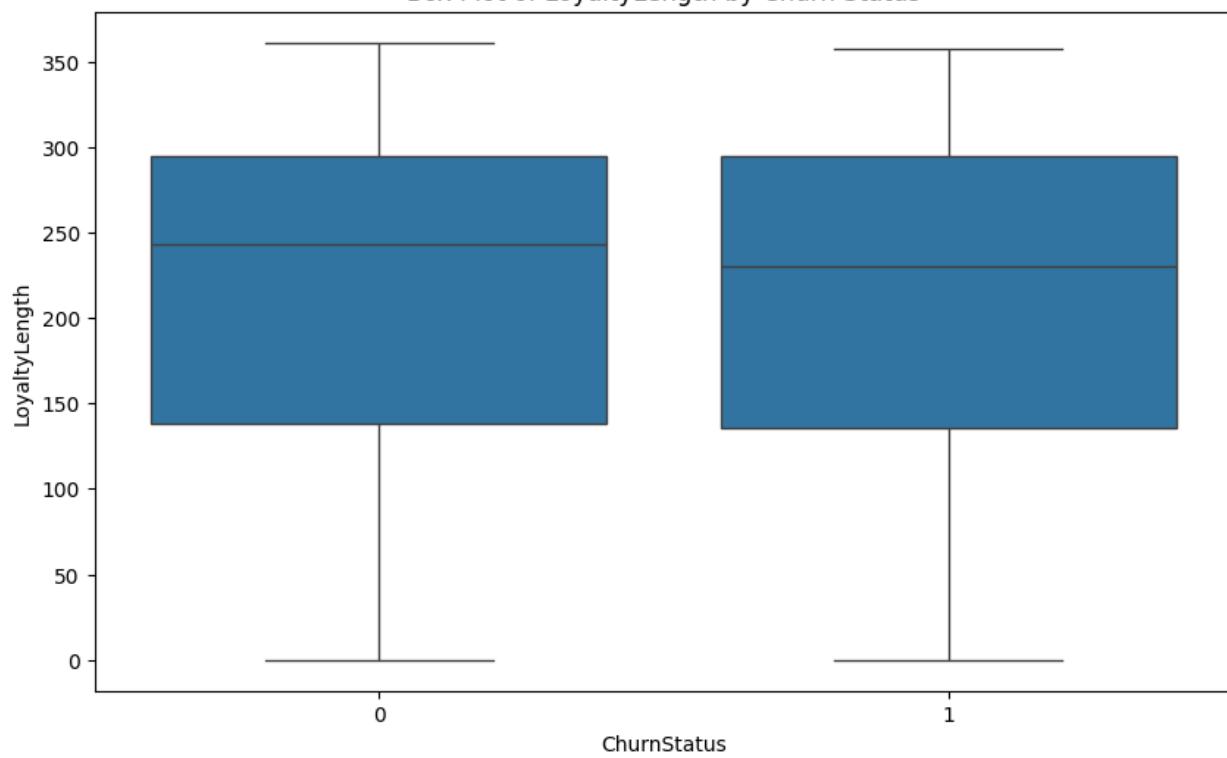
Box Plot of MaxTransaction by Churn Status



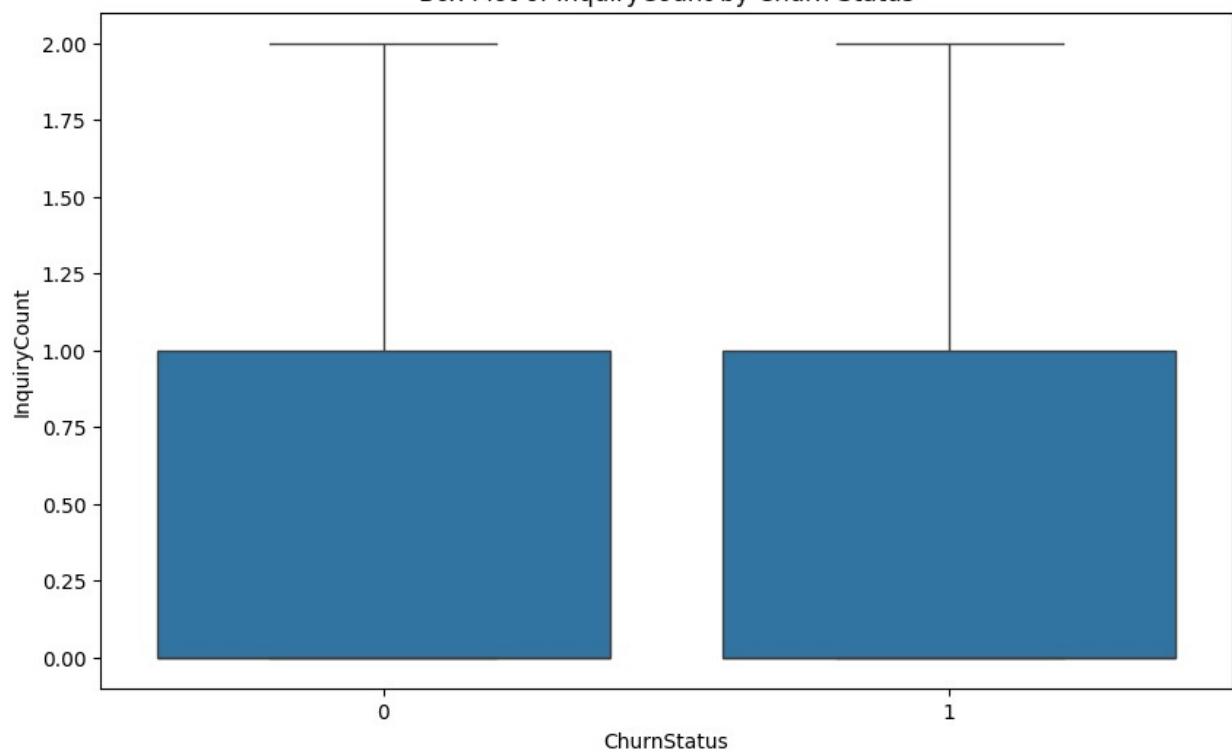
Box Plot of TransactionFrequency by Churn Status



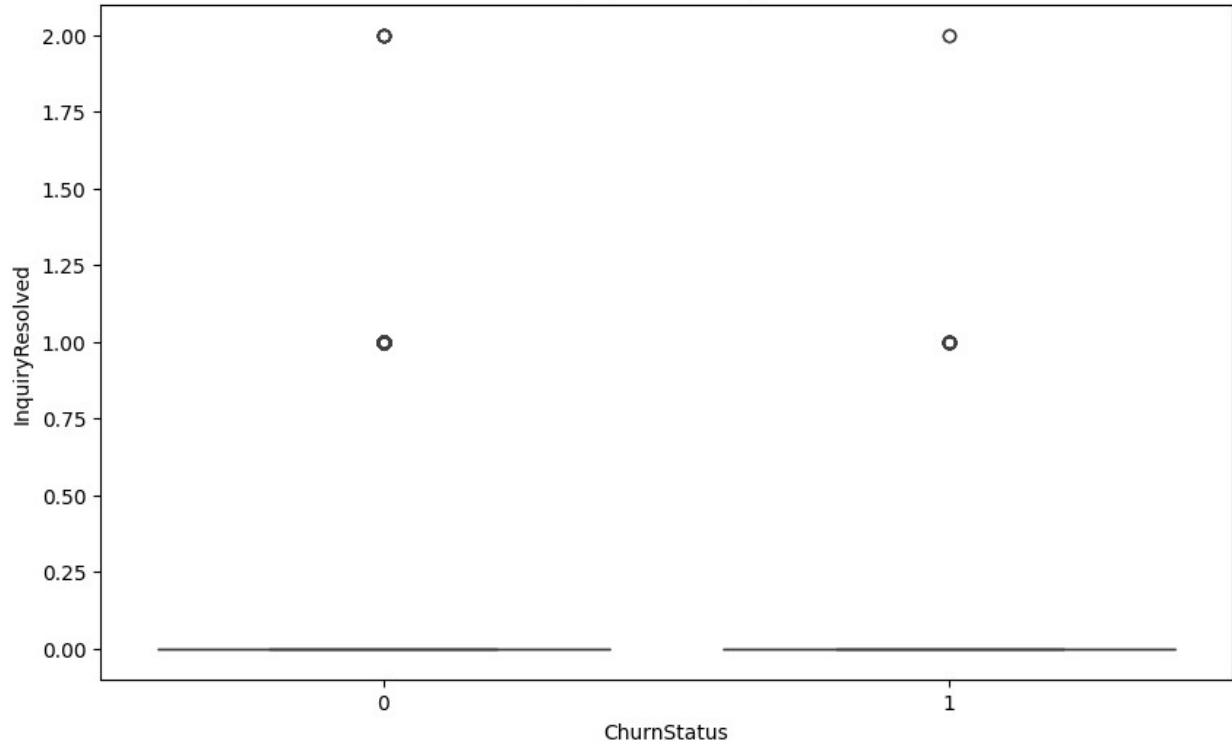
Box Plot of LoyaltyLength by Churn Status



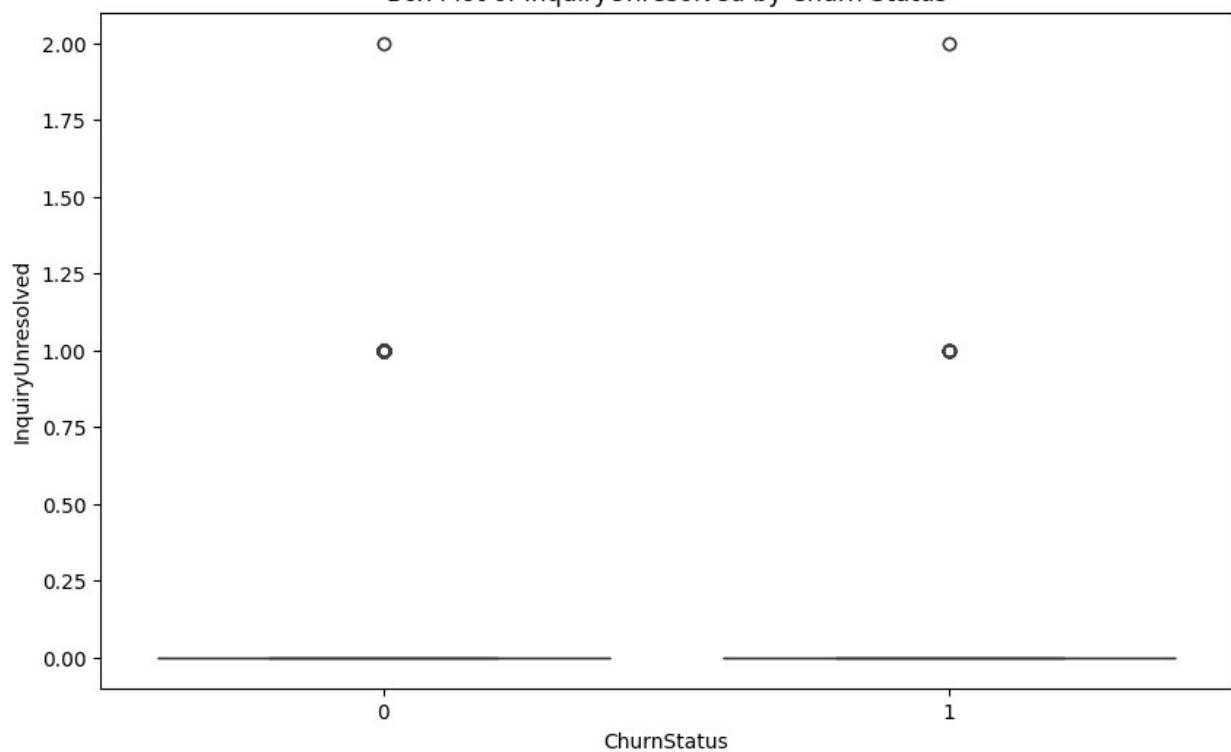
Box Plot of InquiryCount by Churn Status



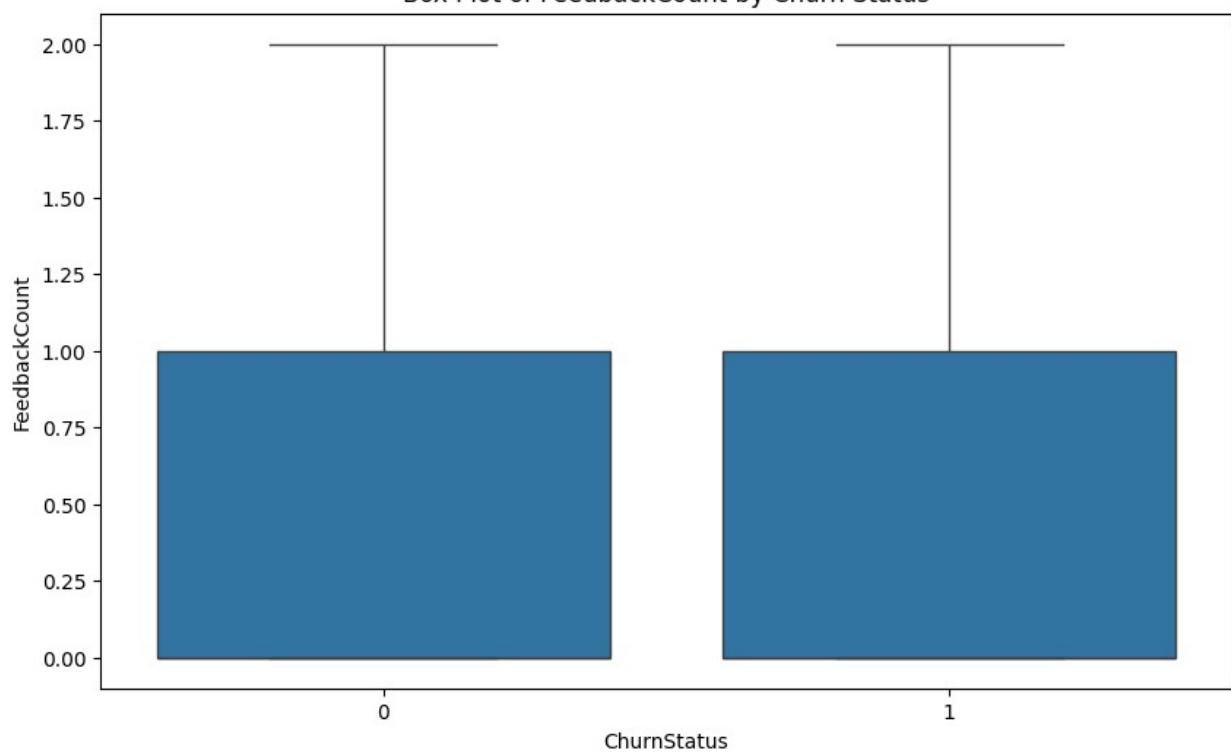
Box Plot of InquiryResolved by Churn Status



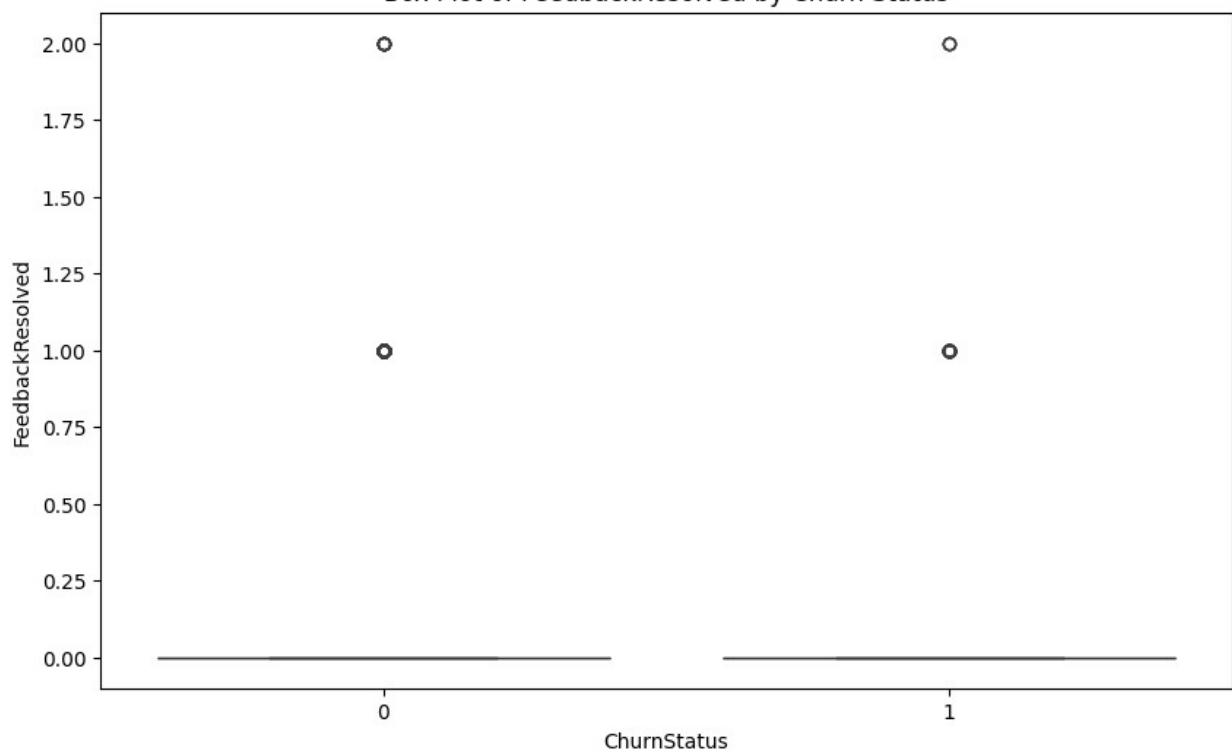
Box Plot of InquiryUnresolved by Churn Status



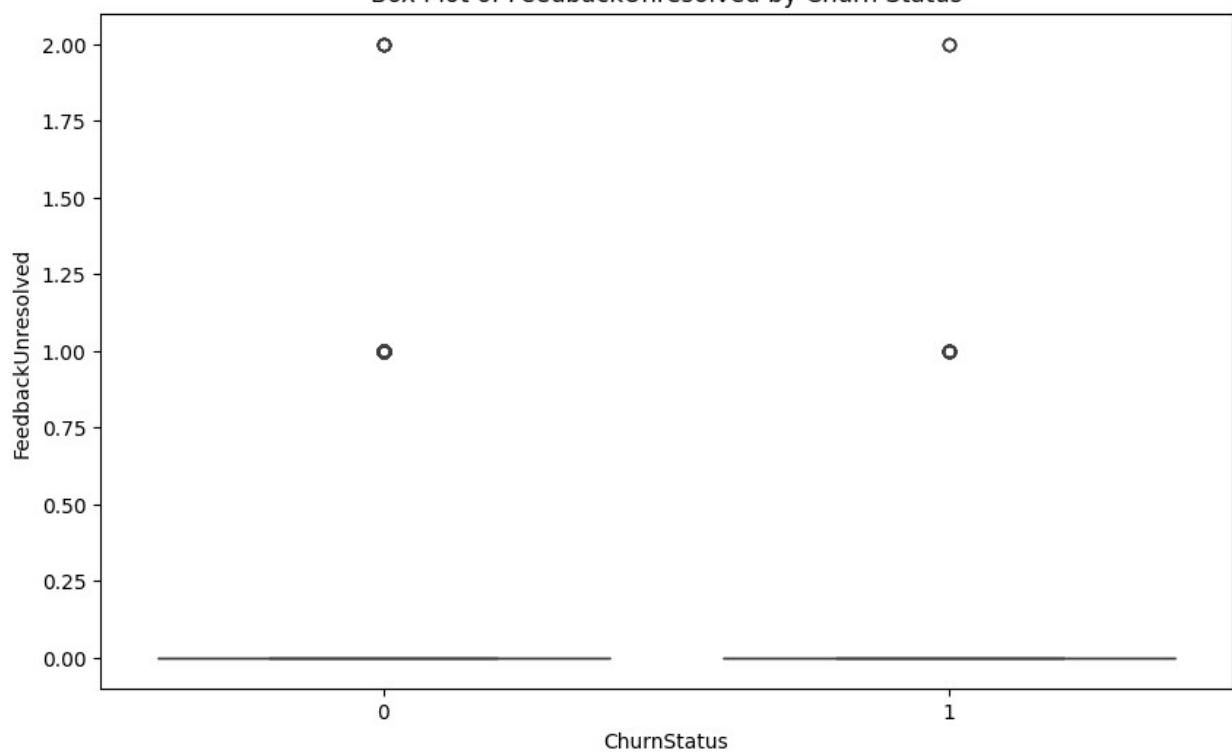
Box Plot of FeedbackCount by Churn Status



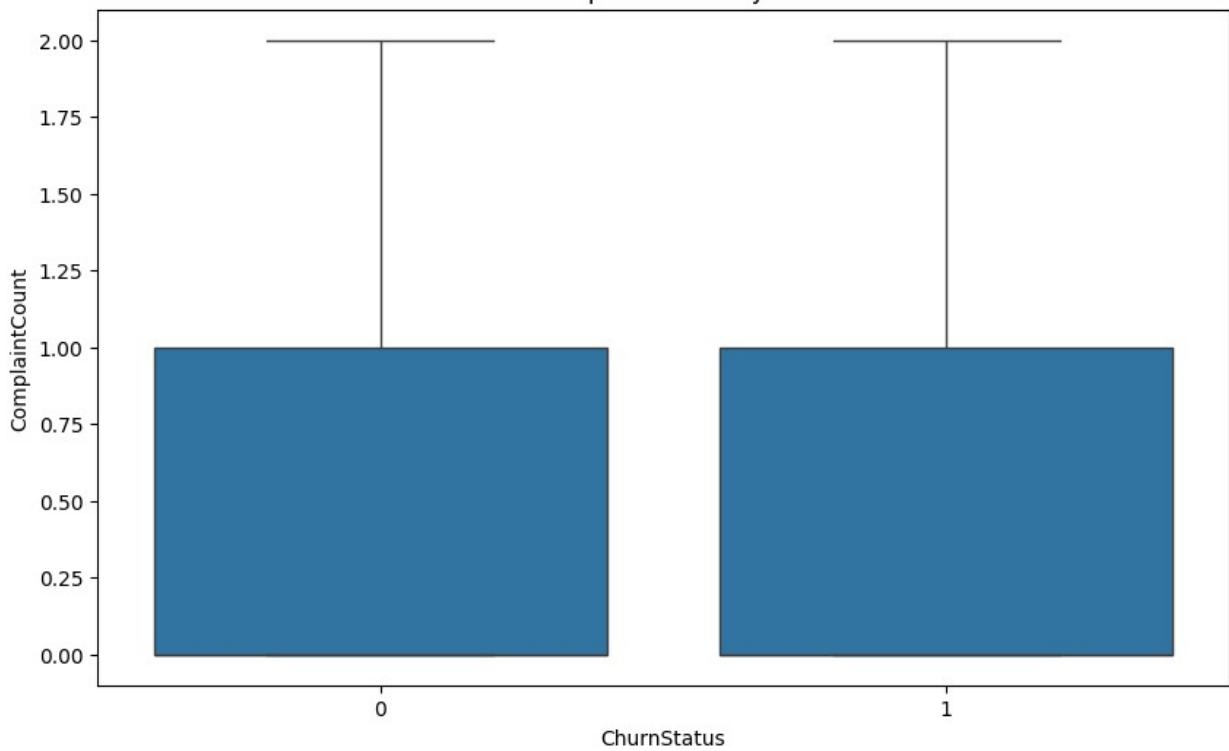
Box Plot of FeedbackResolved by Churn Status



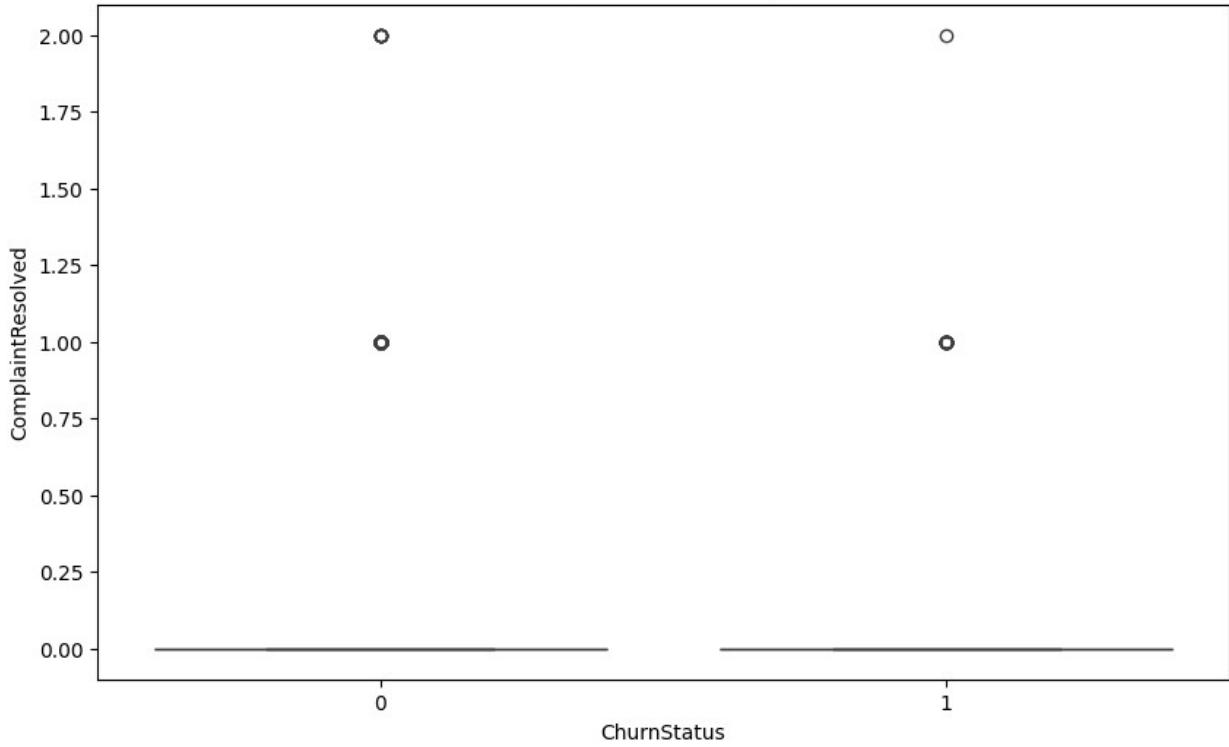
Box Plot of FeedbackUnresolved by Churn Status

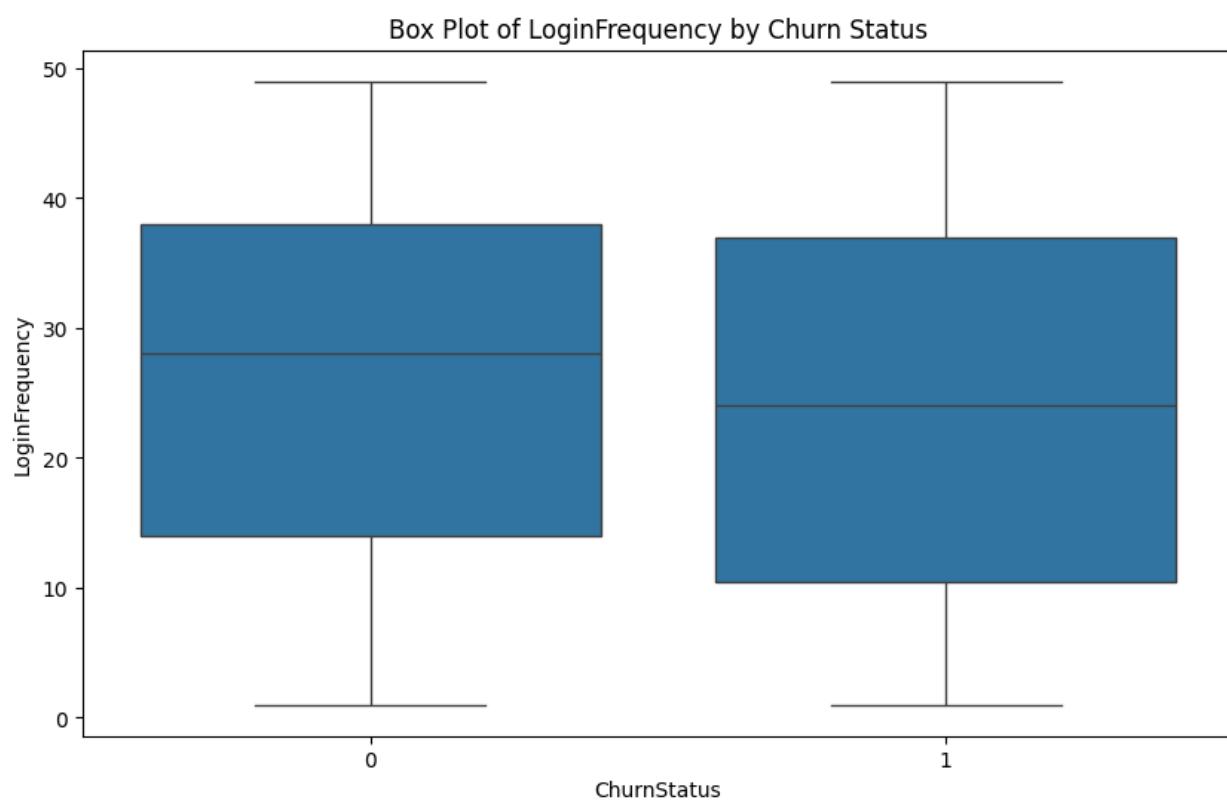
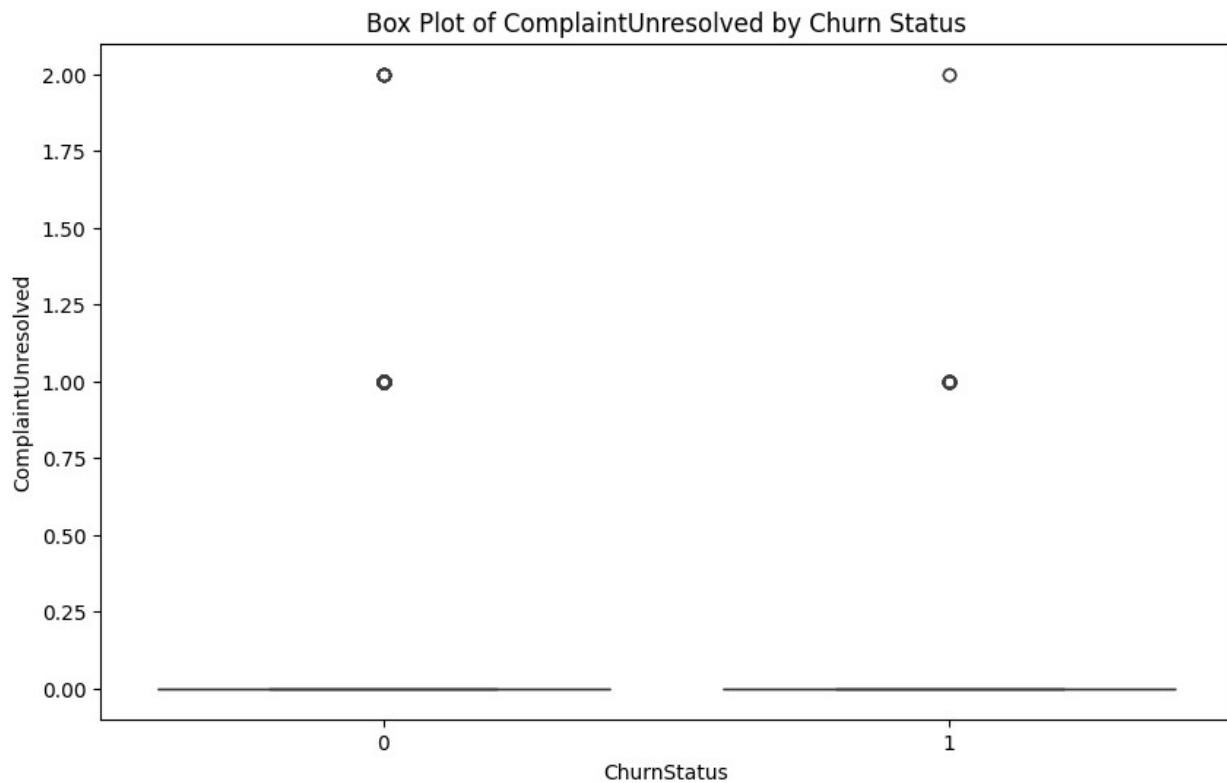


Box Plot of ComplaintCount by Churn Status



Box Plot of ComplaintResolved by Churn Status





Encode categorical data to numbers to prep for algorithm

```
# Encode categorical variables using mapping for the IncomeLevel as order matters.
```

```
Customer_Data['IncomeLevel'] = Customer_Data['IncomeLevel'].map({  
    'Low': 1,  
    'Medium': 2,  
    'High': 3  
})
```

```
display(Customer_Data.head())
```

	CustomerID	Age	Gender	MaritalStatus	IncomeLevel	ChurnStatus
TotalSpent \						
0 416.50	1	62	M	Single	1	0
1 1547.42	2	65	M	Married	1	1
2 1702.98	3	18	M	Single	1	0
3 917.29	4	21	M	Widowed	1	0
4 2001.49	5	21	M	Divorced	2	0

	MinTransaction	MaxTransaction	TransactionFrequency
MostRecentTransaction \			
0 2022-03-27	416.50	416.50	1
1 2022-11-19	54.96	397.37	7
2 2022-10-08	51.07	419.95	6
3 2022-12-27	44.22	382.39	5
4 2022-12-21	69.86	475.69	8

	FirstTransaction	LoyaltyLength	InquiryCount	InquiryResolved	\
0	2022-03-27	0	1.00	1.00	
1	2022-01-09	314	1.00	1.00	
2	2022-02-11	239	1.00	1.00	
3	2022-05-22	219	2.00	1.00	
4	2022-02-21	303	0.00	0.00	

	InquiryUnresolved	FeedbackCount	FeedbackResolved
FeedbackUnresolved \			
0 0.00	0.00	0.00	0.00
1 0.00	0.00	0.00	0.00
0 0.00			

```

2           0.00           0.00           0.00
0.00
3           1.00           0.00           0.00
0.00
4           0.00           0.00           0.00
0.00

    ComplaintCount  ComplaintResolved  ComplaintUnresolved
LastLoginDate \
0           0.00           0.00           0.00      2023-10-
21
1           0.00           0.00           0.00      2023-12-
05
2           0.00           0.00           0.00      2023-11-
15
3           0.00           0.00           0.00      2023-08-
25
4           0.00           0.00           0.00      2023-10-
27

    LoginFrequency ServiceUsage
0           34   Mobile App
1           5    Website
2           3    Website
3           2    Website
4          41    Website

# Encode Gender using one-hot encoding since order does not matter

isMale = pd.get_dummies(Customer_Data['Gender'], drop_first=True)
Customer_Data['Gender'] = isMale
Customer_Data['Gender'] = Customer_Data['Gender'].astype(int) # Change
from true to an integer 1

# Rename Columns for clarity
Customer_Data = Customer_Data.rename(columns={'Gender': 'isMale'})

# Drop the date columns as they have served their purpose
Customer_Data = Customer_Data.drop(columns=['MostRecentTransaction',
'FirstTransaction', 'LastLoginDate'])

display(Customer_Data.head())

    CustomerID  Age  isMale MaritalStatus  IncomeLevel  ChurnStatus \
0           1   62       1      Single        1            0
1           2   65       1     Married        1            1
2           3   18       1      Single        1            0
3           4   21       1     Widowed        1            0
4           5   21       1     Divorced        2            0

    TotalSpent  MinTransaction  MaxTransaction  TransactionFrequency \

```

0	416.50	416.50	416.50	1		
1	1547.42	54.96	397.37	7		
2	1702.98	51.07	419.95	6		
3	917.29	44.22	382.39	5		
4	2001.49	69.86	475.69	8		
0	LoyaltyLength	InquiryCount	InquiryResolved	InquiryUnresolved		
0	0	1.00	1.00	0.00		
1	314	1.00	1.00	0.00		
2	239	1.00	1.00	0.00		
3	219	2.00	1.00	1.00		
4	303	0.00	0.00	0.00		
0	FeedbackCount	FeedbackResolved	FeedbackUnresolved	ComplaintCount		
0	0.00	0.00	0.00	0.00		
1	0.00	0.00	0.00	0.00		
2	0.00	0.00	0.00	0.00		
3	0.00	0.00	0.00	0.00		
4	0.00	0.00	0.00	0.00		
0	ComplaintResolved	ComplaintUnresolved	LoginFrequency	ServiceUsage		
0	0.00	0.00	34	Mobile App		
1	0.00	0.00	5	Website		
2	0.00	0.00	3	Website		
3	0.00	0.00	2	Website		
4	0.00	0.00	41	Website		
# Encode all the other categorical variables using one-hot encoding						
Customer_Data = pd.get_dummies(Customer_Data,						
columns=['MaritalStatus', 'ServiceUsage'], dtype=int)						
display(Customer_Data.head())						
Customer_Data.info()						
0	CustomerID	Age	isMale	IncomeLevel	ChurnStatus	TotalSpent
0	1	62	1	1	0	416.50
1	2	65	1	1	1	1547.42
2	3	18	1	1	0	1702.98
3	4	21	1	1	0	917.29
4	5	21	1	2	0	2001.49

	MinTransaction	MaxTransaction	TransactionFrequency	LoyaltyLength
0	416.50	416.50	1	0
1	54.96	397.37	7	314
2	51.07	419.95	6	239
3	44.22	382.39	5	219
4	69.86	475.69	8	303
	InquiryCount	InquiryResolved	InquiryUnresolved	FeedbackCount
0	1.00	1.00	0.00	0.00
1	1.00	1.00	0.00	0.00
2	1.00	1.00	0.00	0.00
3	2.00	1.00	1.00	0.00
4	0.00	0.00	0.00	0.00
	FeedbackResolved	FeedbackUnresolved	ComplaintCount	ComplaintResolved
0	0.00	0.00	0.00	0.00
1	0.00	0.00	0.00	0.00
2	0.00	0.00	0.00	0.00
3	0.00	0.00	0.00	0.00
4	0.00	0.00	0.00	0.00
	ComplaintUnresolved	LoginFrequency	MaritalStatus_Divorced	MaritalStatus_Married
0	0.00	34	0	0
1	0.00	5	0	1
2	0.00	3	0	0
3	0.00	2	0	0
4	0.00	41	1	0
	MaritalStatus_Married	MaritalStatus_Single	MaritalStatus_Widowed	MaritalStatus_Divorced
0	0	1	0	0
1	1	0	0	0
2	0	1	0	0
3	0	0	1	0

```
4          0          0          0
```

	ServiceUsage_Mobile App	ServiceUsage_Online Banking
ServiceUsage_Website		
0	1	0
0		
1	0	0
1		
2	0	0
1		
3	0	0
1		
4	0	0
1		

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 27 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   CustomerID      1000 non-null    int64  
 1   Age              1000 non-null    int64  
 2   isMale           1000 non-null    int64  
 3   IncomeLevel     1000 non-null    int64  
 4   ChurnStatus     1000 non-null    int64  
 5   TotalSpent      1000 non-null    float64 
 6   MinTransaction  1000 non-null    float64 
 7   MaxTransaction  1000 non-null    float64 
 8   TransactionFrequency  1000 non-null    int64  
 9   LoyaltyLength   1000 non-null    int64  
 10  InquiryCount   1000 non-null    float64 
 11  InquiryResolved 1000 non-null    float64 
 12  InquiryUnresolved 1000 non-null    float64 
 13  FeedbackCount  1000 non-null    float64 
 14  FeedbackResolved 1000 non-null    float64 
 15  FeedbackUnresolved 1000 non-null    float64 
 16  ComplaintCount 1000 non-null    float64 
 17  ComplaintResolved 1000 non-null    float64 
 18  ComplaintUnresolved 1000 non-null    float64 
 19  LoginFrequency  1000 non-null    int64  
 20  MaritalStatus_Divorced 1000 non-null    int64  
 21  MaritalStatus_Married 1000 non-null    int64  
 22  MaritalStatus_Single 1000 non-null    int64  
 23  MaritalStatus_Widowed 1000 non-null    int64  
 24  ServiceUsage_Mobile App 1000 non-null    int64  
 25  ServiceUsage_Online Banking 1000 non-null    int64  
 26  ServiceUsage_Website 1000 non-null    int64  
dtypes: float64(12), int64(15)
memory usage: 211.1 KB
```

```

# Export to excel for submission

Customer_Data.to_excel('Customer_Data_Cleaned.xlsx', index=False)

# Used StandardScaler() to calculate z scores so all numeric values
# hold the same weight.
from sklearn.preprocessing import StandardScaler

# Identify columns to scale drop non-numeric and target variable as
# well as CustomerID
columns_to_scale =
Customer_Data.select_dtypes(include='number').columns.drop(['CustomerI
D', 'ChurnStatus'])

# Scale them
scaler = StandardScaler()
scaled_values = scaler.fit_transform(Customer_Data[columns_to_scale])
Customer_Data[columns_to_scale] =
scaler.fit_transform(Customer_Data[columns_to_scale])

display(Customer_Data.head())

```

	CustomerID	Age	isMale	IncomeLevel	ChurnStatus	TotalSpent
0	1	1.23	1.03	-1.25	0	-1.15
1	2	1.43	1.03	-1.25	1	0.38
2	3	-1.66	1.03	-1.25	0	0.59
3	4	-1.46	1.03	-1.25	0	-0.47
4	5	-1.46	1.03	-0.03	0	0.99

	MinTransaction	MaxTransaction	TransactionFrequency	LoyaltyLength
0	3.06	0.25	-1.56	-1.89
1	-0.52	0.07	0.75	0.96
2	-0.55	0.28	0.36	0.28
3	-0.62	-0.07	-0.02	0.10
4	-0.37	0.80	1.13	0.86

	InquiryCount	InquiryResolved	InquiryUnresolved	FeedbackCount
0	1.34	2.09	-0.39	-0.62
1	1.34	2.09	-0.39	-0.62
2	1.34	2.09	-0.39	-0.62
3	3.27	2.09	2.41	-0.62
4	-0.59	-0.42	-0.39	-0.62

	FeedbackResolved	FeedbackUnresolved	ComplaintCount
--	------------------	--------------------	----------------

```

ComplaintResolved \
0           -0.47          -0.41          -0.61          -
0.41
1           -0.47          -0.41          -0.61          -
0.41
2           -0.47          -0.41          -0.61          -
0.41
3           -0.47          -0.41          -0.61          -
0.41
4           -0.47          -0.41          -0.61          -
0.41

    ComplaintUnresolved  LoginFrequency  MaritalStatus_Divorced \
0           -0.44           0.58          -0.57
1           -0.44          -1.49          -0.57
2           -0.44          -1.63          -0.57
3           -0.44          -1.70          -0.57
4           -0.44           1.07           1.74

    MaritalStatus_Married  MaritalStatus_Single  MaritalStatus_Widowed \
\0           -0.59           1.91          -0.62
1            1.68          -0.52          -0.62
2           -0.59           1.91          -0.62
3           -0.59          -0.52           1.62
4           -0.59          -0.52          -0.62

    ServiceUsage_Mobile App  ServiceUsage_Online Banking
ServiceUsage_Website
0                  1.39          -0.73
-0.67
1                 -0.72          -0.73
1.50
2                 -0.72          -0.73
1.50
3                 -0.72          -0.73
1.50
4                 -0.72          -0.73
1.50

# Check Correlation of numerical features with Churn Status

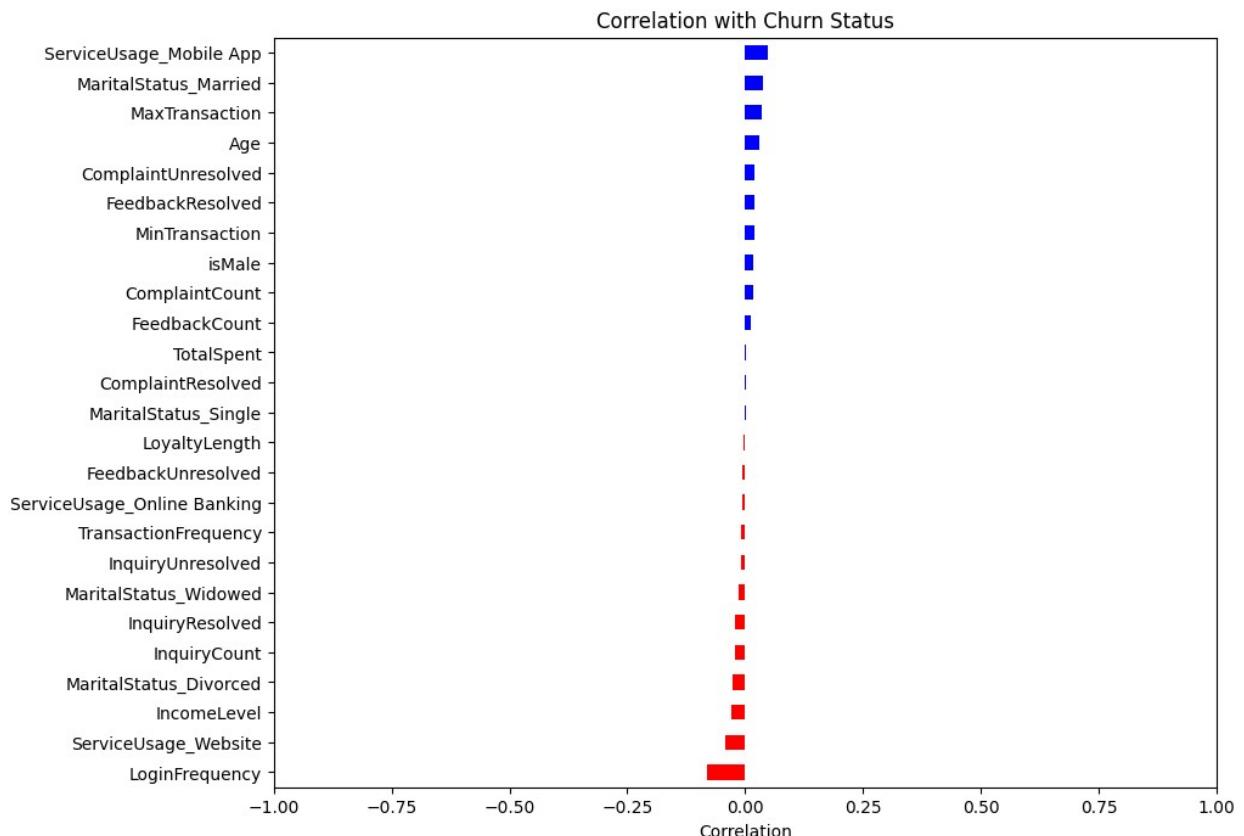
churn_corr = Customer_Data.select_dtypes(include='number').corr()
['ChurnStatus'].drop(['ChurnStatus', 'CustomerID']).sort_values() # include all numerical values -ChurnStatus and CustomerID

```

```

plt.figure(figsize=(10, 8))
churn_corr.plot(kind='barh', color=['red' if x < 0 else 'blue' for x
in churn_corr]) # a horizontal bar plot with red for negative
correlation and blue for positive correlation
plt.title('Correlation with Churn Status')
plt.xlabel('Correlation')
plt.xlim(-1, 1)
plt.show()

```



Observation:

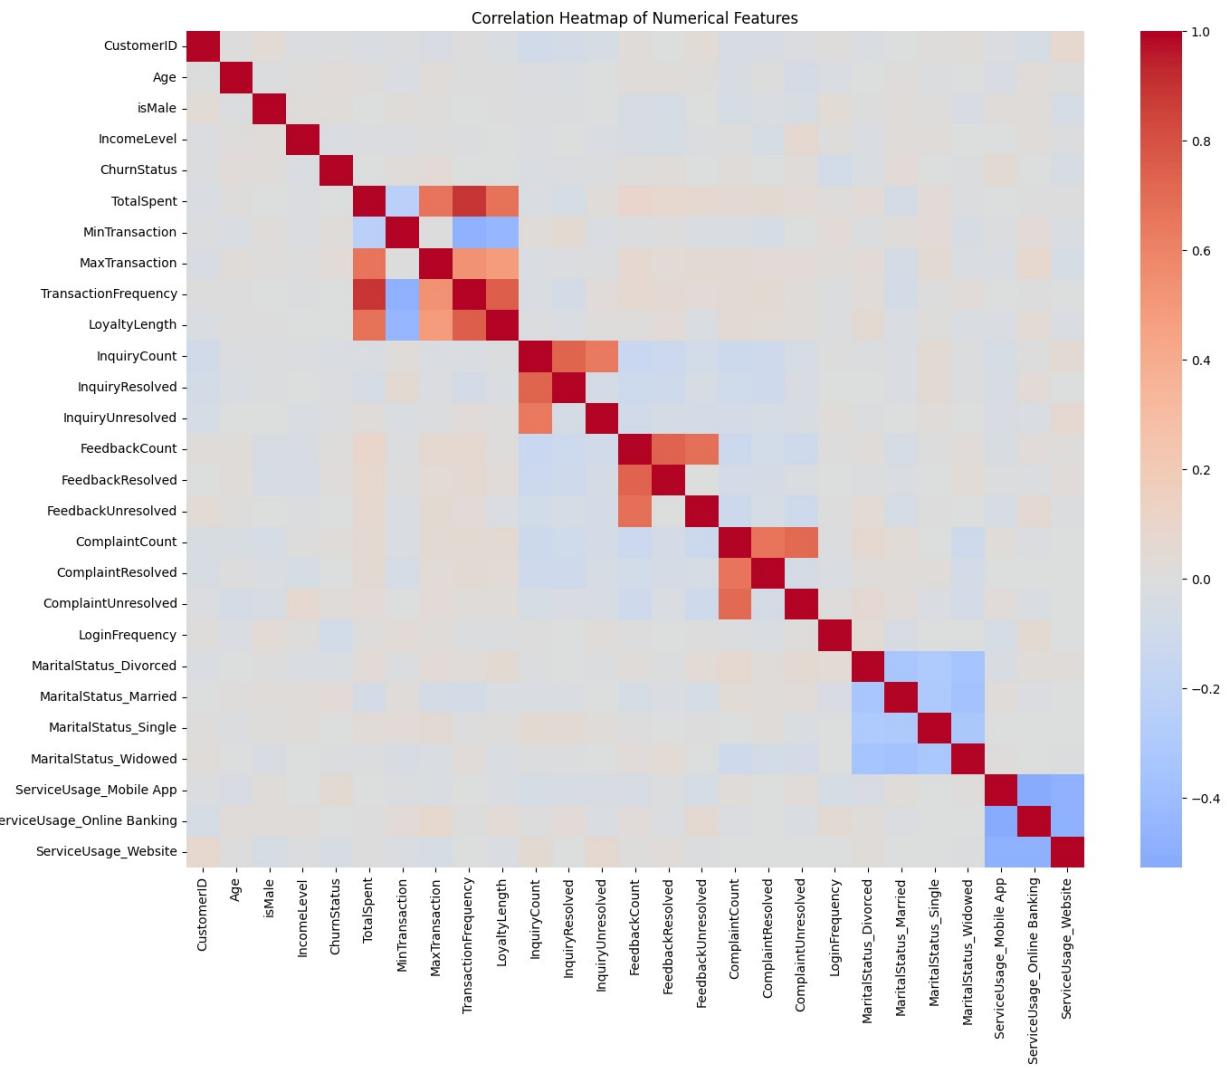
There is little to no correlation with any of the features to ChurnStatus the strongest correlation is LoginFrequency at approximately 0.08. It might as well be nothing.

```

# do a correlation heatmap to see how features correlate with each
other
plt.figure(figsize=(16, 12))
correlation_matrix =
Customer_Data.select_dtypes(include='number').corr() # select only
numerical columns
sns.heatmap(correlation_matrix, annot=False, cmap='coolwarm',
center=0)

```

```
plt.title('Correlation Heatmap of Numerical Features')
plt.show()
```



Observation

There is little to no correlation with any of the features with each other.

Why It Matters: This indicates a systemic issue with data collection that must be addressed. The data we collect is next to useless for driving any meaningful business decision.

```
display(Customer_Data.head())
```

	CustomerID	Age	isMale	IncomeLevel	ChurnStatus	TotalSpent	\
0	1	1.23	1.03		-1.25	0	-1.15
1	2	1.43	1.03		-1.25	1	0.38
2	3	-1.66	1.03		-1.25	0	0.59
3	4	-1.46	1.03		-1.25	0	-0.47
4	5	-1.46	1.03		-0.03	0	0.99

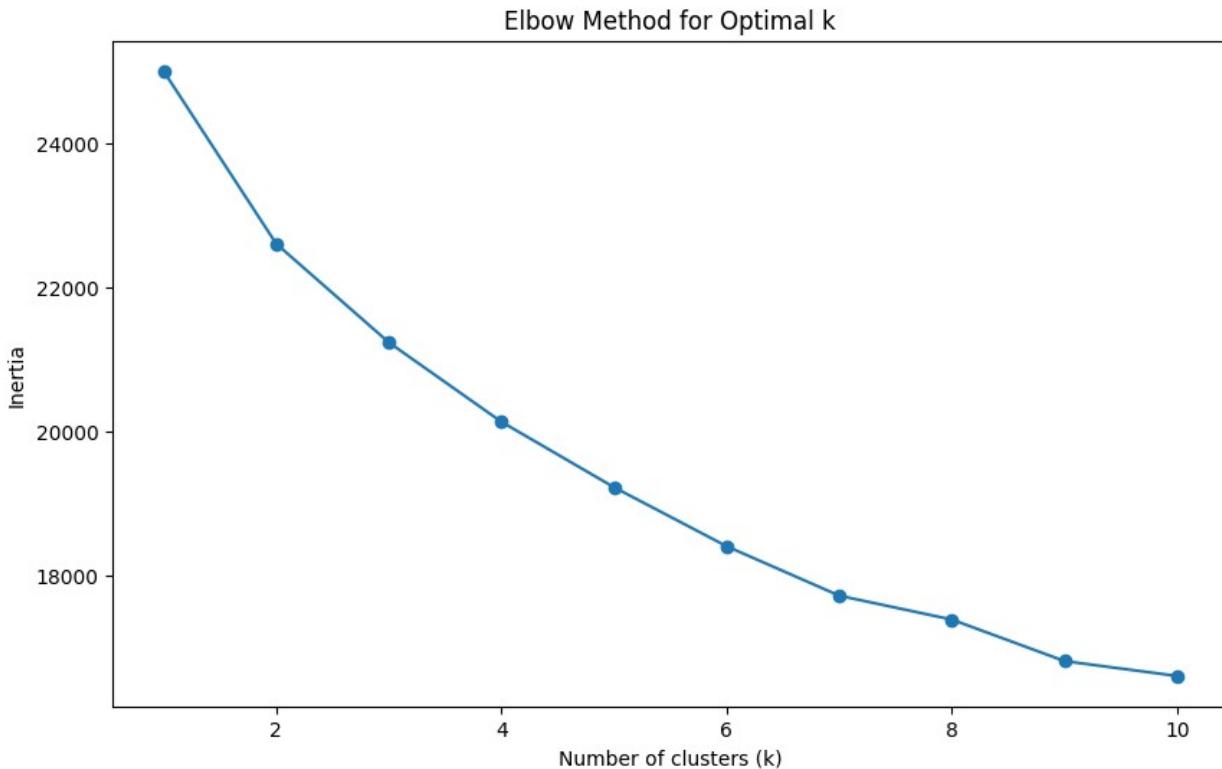
	MinTransaction	MaxTransaction	TransactionFrequency	LoyaltyLength
0	3.06	0.25	-1.56	-1.89
1	-0.52	0.07	0.75	0.96
2	-0.55	0.28	0.36	0.28
3	-0.62	-0.07	-0.02	0.10
4	-0.37	0.80	1.13	0.86
	InquiryCount	InquiryResolved	InquiryUnresolved	FeedbackCount
0	1.34	2.09	-0.39	-0.62
1	1.34	2.09	-0.39	-0.62
2	1.34	2.09	-0.39	-0.62
3	3.27	2.09	2.41	-0.62
4	-0.59	-0.42	-0.39	-0.62
	FeedbackResolved	FeedbackUnresolved	ComplaintCount	
ComplaintResolved	-0.47	-0.41	-0.61	-
0	-0.47	-0.41	-0.61	-
0.41	-0.47	-0.41	-0.61	-
1	-0.47	-0.41	-0.61	-
0.41	-0.47	-0.41	-0.61	-
2	-0.47	-0.41	-0.61	-
0.41	-0.47	-0.41	-0.61	-
3	-0.47	-0.41	-0.61	-
0.41	-0.47	-0.41	-0.61	-
4	-0.47	-0.41	-0.61	-
0.41	-0.47	-0.41	-0.61	-
	ComplaintUnresolved	LoginFrequency	MaritalStatus_Divorced	
0	-0.44	0.58	-0.57	-
1	-0.44	-1.49	-0.57	-
2	-0.44	-1.63	-0.57	-
3	-0.44	-1.70	-0.57	-
4	-0.44	1.07	1.74	-
	MaritalStatus_Married	MaritalStatus_Single	MaritalStatus_Widowed	
\	-0.59	1.91	-0.62	-
0	1.68	-0.52	-0.62	-
1	-0.59	1.91	-0.62	-
2	-0.59	-0.52	1.62	-
3	-0.59	-0.52	-0.62	-

	-0.59	-0.52	-0.62
--	-------	-------	-------

ServiceUsage_Mobile App	ServiceUsage_Online Banking	
ServiceUsage_Website		
0	1.39	-0.73
-0.67		
1	-0.72	-0.73
1.50		
2	-0.72	-0.73
1.50		
3	-0.72	-0.73
1.50		
4	-0.72	-0.73
1.50		

Cluster Analysis to identify customer segments to see if we can identify any patterns among churners vs non-churners

```
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
# Select features for clustering (excluding CustomerID and ChurnStatus)
features =
Customer_Data.select_dtypes(include='number').drop(columns=['CustomerID', 'ChurnStatus'])
# Determine optimal number of clusters using Elbow Method
inertia = []
for k in range(1, 11):
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
    kmeans.fit(features)
    inertia.append(kmeans.inertia_)
plt.figure(figsize=(10, 6))
plt.plot(range(1, 11), inertia, marker='o')
plt.title('Elbow Method for Optimal k')
plt.xlabel('Number of clusters (k)')
plt.ylabel('Inertia')
plt.show()
```



```

# Create combined unresolved interactions column
Customer_Data['TotalUnresolved'] =
(Customer_Data['FeedbackUnresolved'] +
Customer_Data['ComplaintUnresolved'] +
Customer_Data['InquiryUnresolved'])

# Now run clustering and aggregate
kmeans = KMeans(n_clusters=9, random_state=42, n_init=10)
Customer_Data['Cluster'] = kmeans.fit_predict(features)

cluster_summary = Customer_Data.groupby('Cluster').agg(
    Count = ('CustomerID', 'count'),
    Age = ('Age', 'mean'),
    TotalUnresolved = ('TotalUnresolved', 'mean'),
    TotalSpent = ('TotalSpent', 'mean'),
    TransactionFrequency = ('TransactionFrequency', 'mean'),
    LoyaltyLength = ('LoyaltyLength', 'mean'),
    LoginFrequency = ('LoginFrequency', 'mean'),
    ChurnRate = ('ChurnStatus', 'mean')
).reset_index()

display(cluster_summary)

```

	Cluster	Count	Age	TotalUnresolved	TotalSpent
TransactionFrequency \	0	115	0.01	-0.50	0.22
0.25	1	118	-0.07	-0.50	0.10
0.10	2	108	-0.01	-0.41	-1.38
1.28	3	234	0.07	-1.23	0.32
0.41	4	92	-0.16	1.79	0.32
0.23	5	79	-0.19	-0.31	-1.06
1.43	6	78	0.12	1.63	0.44
0.35	7	114	0.02	2.27	0.30
0.36	8	62	0.19	-1.23	0.36
0.41					
	LoyaltyLength	LoginFrequency	ChurnRate		
0	0.26	-0.05	0.22		
1	0.24	-0.04	0.19		
2	-1.30	0.01	0.16		
3	0.45	-0.06	0.19		
4	0.30	0.11	0.24		
5	-1.62	0.06	0.20		
6	0.23	0.11	0.21		
7	0.35	0.03	0.20		
8	0.31	-0.02	0.27		

Observations: Group 6 and 8 Churns at a 21% and 27% respectively.

Why It Matters: These are number 1 and number 2 biggest spenders on average. The reason they churn is undetermined because the features have such low correlation. Further investigation is required.

```
# Compare the high-churn vs low-churn clusters
high_churn = Customer_Data[Customer_Data['Cluster'] == 3].mean()
low_churn = Customer_Data[Customer_Data['Cluster'] == 4].mean()

comparison = pd.DataFrame({
    'High Churn (Cluster 3)': high_churn.drop('CustomerID'),
    'Low Churn (Cluster 4)': low_churn.drop('CustomerID'),
    'Difference': high_churn - low_churn
}).round(2)

print(comparison.drop('CustomerID'))
```

	High Churn (Cluster 3)	Low Churn
(Cluster 4) \		
Age	0.07	-
0.16		
ChurnStatus	0.19	
0.24		
Cluster	3.00	
4.00		
ComplaintCount	-0.61	
1.44		
ComplaintResolved	-0.41	-
0.41		
ComplaintUnresolved	-0.44	
2.29		
FeedbackCount	-0.62	
0.05		
FeedbackResolved	-0.47	
0.17		
FeedbackUnresolved	-0.41	-
0.11		
IncomeLevel	0.05	
0.16		
InquiryCount	-0.59	-
0.59		
InquiryResolved	-0.42	-
0.42		
InquiryUnresolved	-0.39	-
0.39		
LoginFrequency	-0.06	
0.11		
LoyaltyLength	0.45	
0.30		
MaritalStatus_Divorced	-0.00	
0.13		
MaritalStatus_Married	-0.04	
0.10		
MaritalStatus_Single	-0.00	-
0.05		
MaritalStatus_Widowed	0.04	-
0.18		
MaxTransaction	0.35	
0.29		
MinTransaction	-0.33	-
0.01		
ServiceUsage_Mobile App	0.09	
0.13		
ServiceUsage_Online Banking	-0.07	
0.04		
ServiceUsage_Website	-0.02	-
0.17		

TotalSpent	0.32
0.32	
TotalUnresolved	-1.23
1.79	
TransactionFrequency	0.41
0.23	
isMale	0.10
0.02	-
	Difference
Age	0.23
ChurnStatus	-0.05
Cluster	-1.00
ComplaintCount	-2.05
ComplaintResolved	0.00
ComplaintUnresolved	-2.72
FeedbackCount	-0.67
FeedbackResolved	-0.64
FeedbackUnresolved	-0.30
IncomeLevel	-0.11
InquiryCount	-0.00
InquiryResolved	-0.00
InquiryUnresolved	0.00
LoginFrequency	-0.17
LoyaltyLength	0.16
MaritalStatus_Divorced	-0.13
MaritalStatus_Married	-0.14
MaritalStatus_Single	0.04
MaritalStatus_Widowed	0.22
MaxTransaction	0.06
MinTransaction	-0.32
ServiceUsage_Mobile App	-0.04
ServiceUsage_Online Banking	-0.11
ServiceUsage_Website	0.15
TotalSpent	0.00
TotalUnresolved	-3.03
TransactionFrequency	0.18
isMale	0.12